SPOT Tutorial Slides

2021 ECP Annual Meeting

Matthew LeGendre

April 12, 2021

LLNL-PRES-XXXXXX

This work was performed under the auspices of the U.S. Department of Energy by Lawrence Livermore National Laboratory under contract DE-AC52-07NA27344. Lawrence Livermore National Security, LLC

Lawrence Livermore National Laboratory

SPOT: Software Performance and Analysis Tracking

- SPOT is visualization for **collections of performance profiles**. Common uses cases are:
  - Performance comparison of nightly tests. Look for performance regressions.
  - Performance tracking of developer changes. Run an MPI scaling study.
  - Collect performance profiles from users. Understand how users run and hit performance issues.

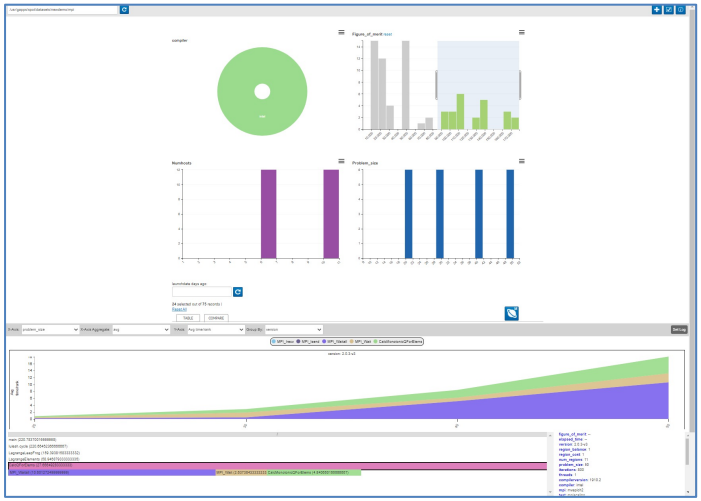

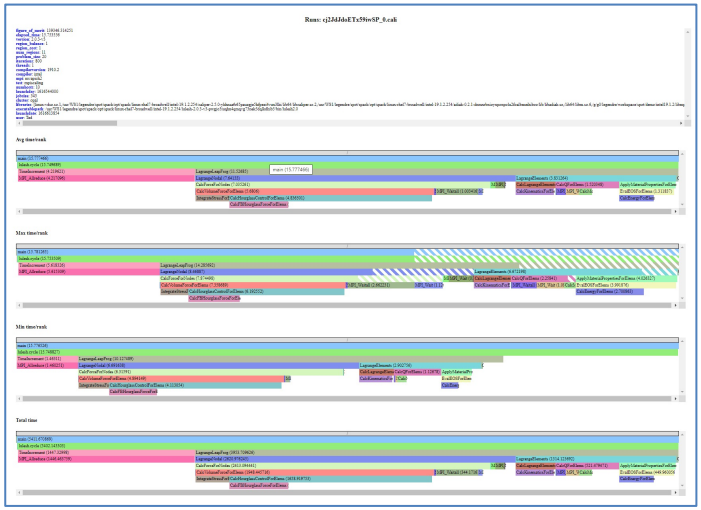

SPOT's Workflow

1. Point SPOT webpage at directory containing *.cali files.

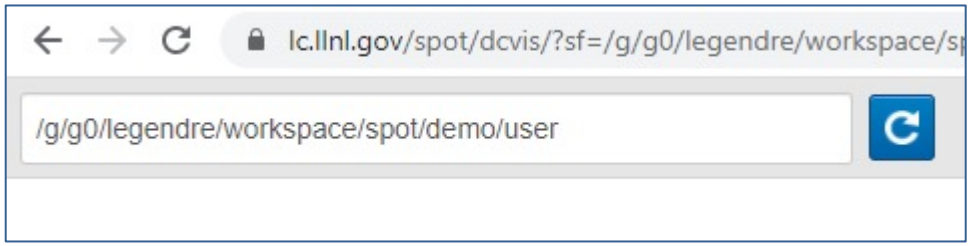

2. Select "interesting" sets of runs to analyze
   - Don't compare performance of 1d test run. vs 3d multi-physics run.

3. Run common analysis with SPOT.
   - i.e., Graph history of performance changes
   - i.e., View performance context trees

4. Run advanced analysis with Hatchet.

SPOT's Landing Page

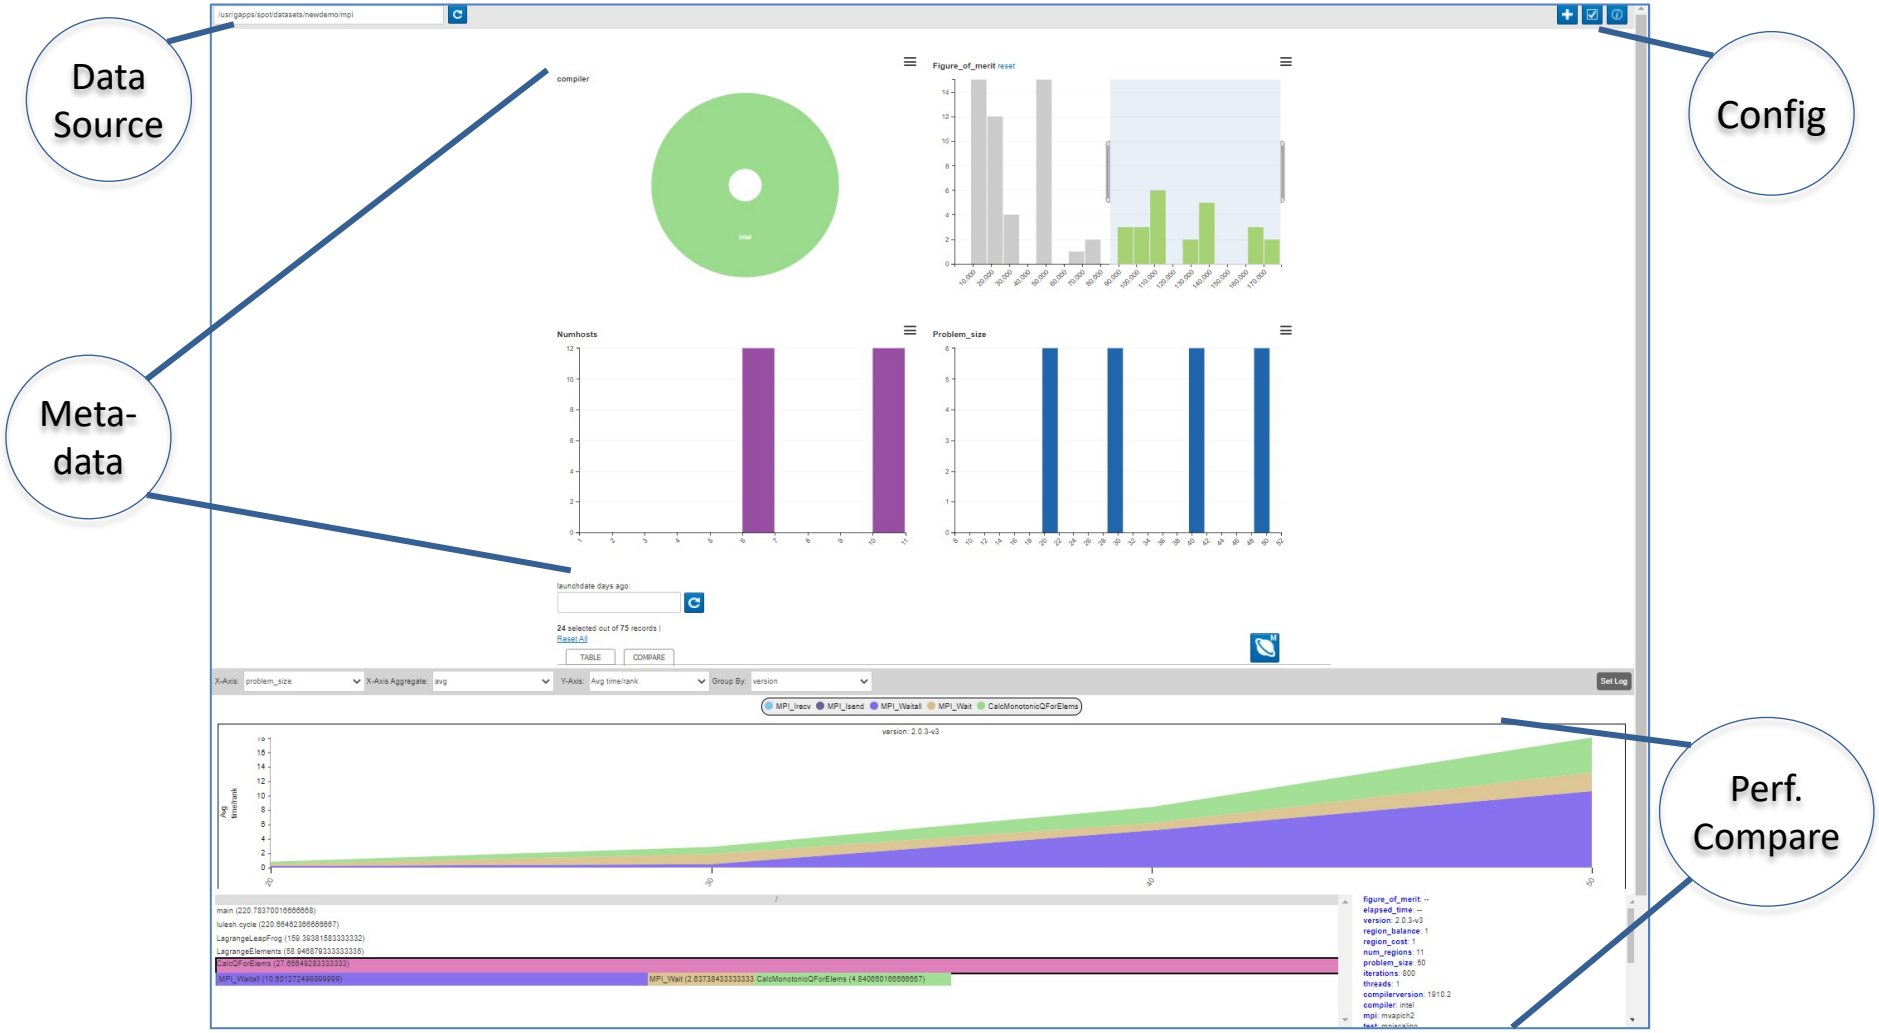

SPOT's Landing Page

Table View

SPOT is Based on CrossFilter

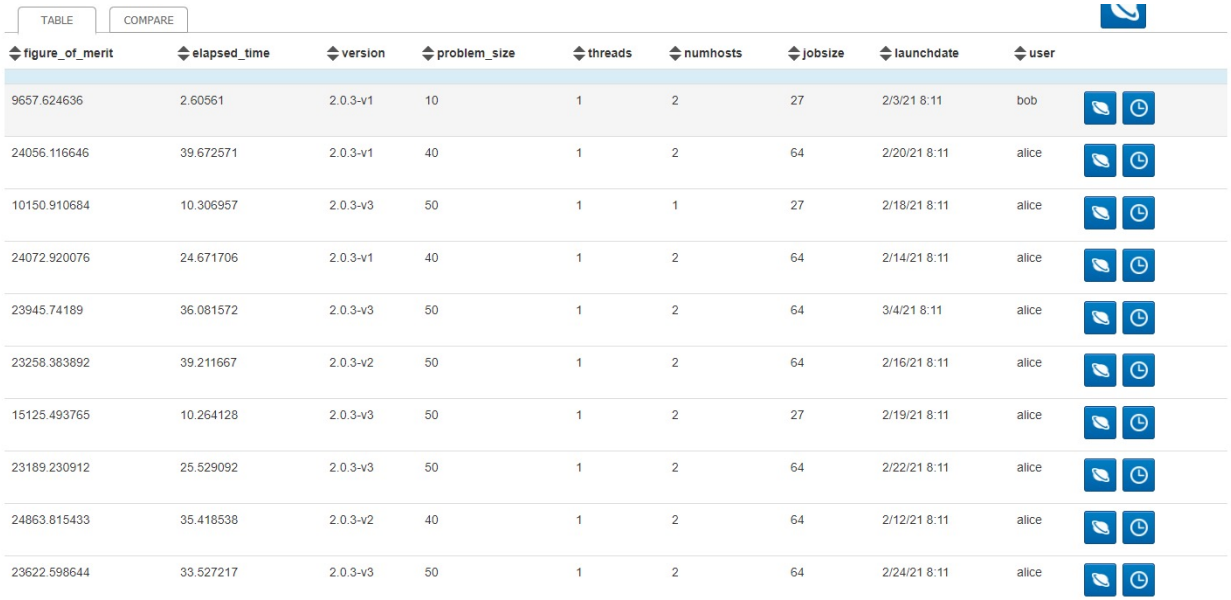

TABLE COMPARE

figure_of_merit	elapsed_time	version	problem_size	threads	numhosts	jobsize	launchdate	user
9657.624636	2.60561	2.0.3-v1	10	1	2	27	2/3/21 8:11	bob
24056.116646	39.672571	2.0.3-v1	40	1	2	64	2/20/21 8:11	alice
10150.910684	10.306957	2.0.3-v3	50	1	1	27	2/18/21 8:11	alice
24072.920076	24.671706	2.0.3-v1	40	1	2	64	2/14/21 8:11	alice
23945.74189	36.081572	2.0.3-v3	50	1	2	64	3/4/21 8:11	alice
23258.383892	39.211667	2.0.3-v2	50	1	2	64	2/16/21 8:11	alice
15125.493765	10.264128	2.0.3-v3	50	1	2	27	2/19/21 8:11	alice
23189.230912	25.529092	2.0.3-v3	50	1	2	64	2/22/21 8:11	alice
24863.815433	35.418538	2.0.3-v2	40	1	2	64	2/12/21 8:11	alice
23622.598644	33.527217	2.0.3-v3	50	1	2	64	2/24/21 8:11	alice

SPOT's Crossfilter: https://dc-js.github.io/dc.js/
D3-Based Crossfilter: https://square.github.io/crossfilter/

- CrossFilter is good at finding correlations in data.
- Crossfilter helps you manage/slice your data to make interesting comparisons.

Each App Run has Two Types of Data

- **Metadata**
  - Name/value data about job.
    - User, host, FOM, input parameters, …
  - Displayed as histograms

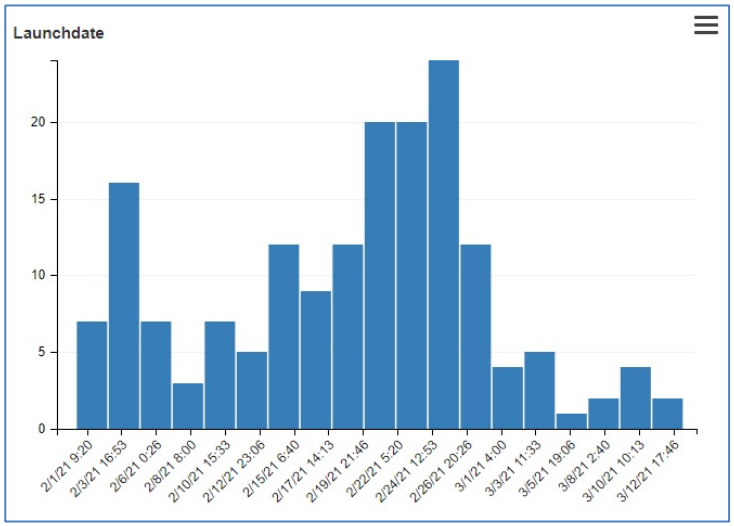

- **Performance Data**
  - Can be different metrics.
    - Avg Time/Rank, Max Time/Rank, bandwidth, …
  - Displayed as a flame graph or timeseries.

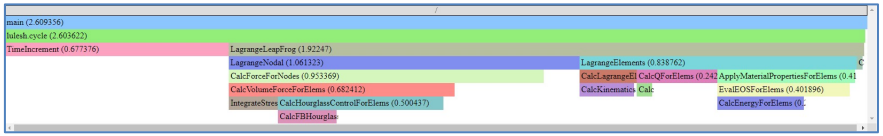

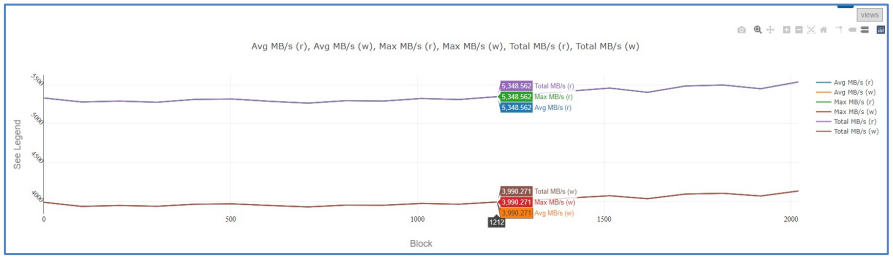

Metadata Can be Filtered

- Select regions of metadata histograms to filter.
  - Filtering in one histogram updates all histograms range.
    - E.g., Filtering Launchdate to ‘Feb 19-Feb26’ will update the ‘Users’ histogram to only show users who ran between those dates.
  - While filtered, comparisons and Jupyter notebooks will operate on the filtered data.

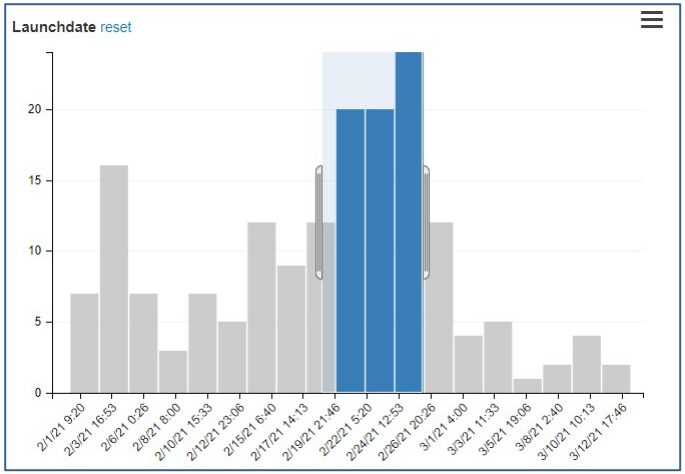

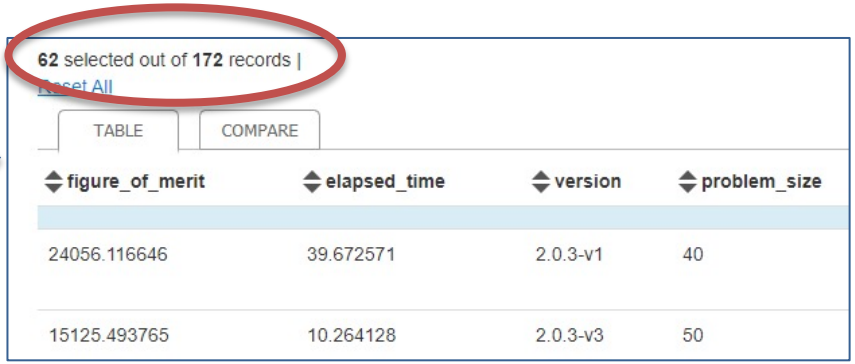

Metadata is Configurable

- Can show/hide metadata:

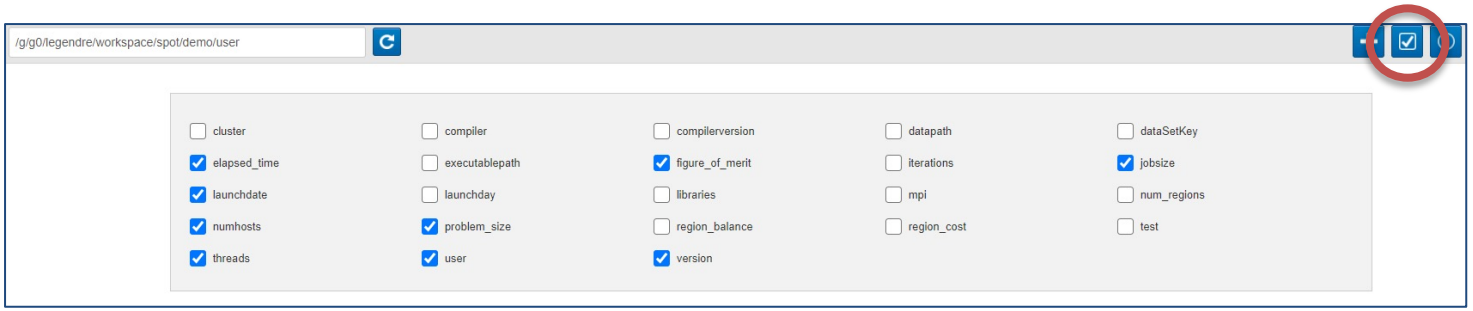

- Can create composite metadata:

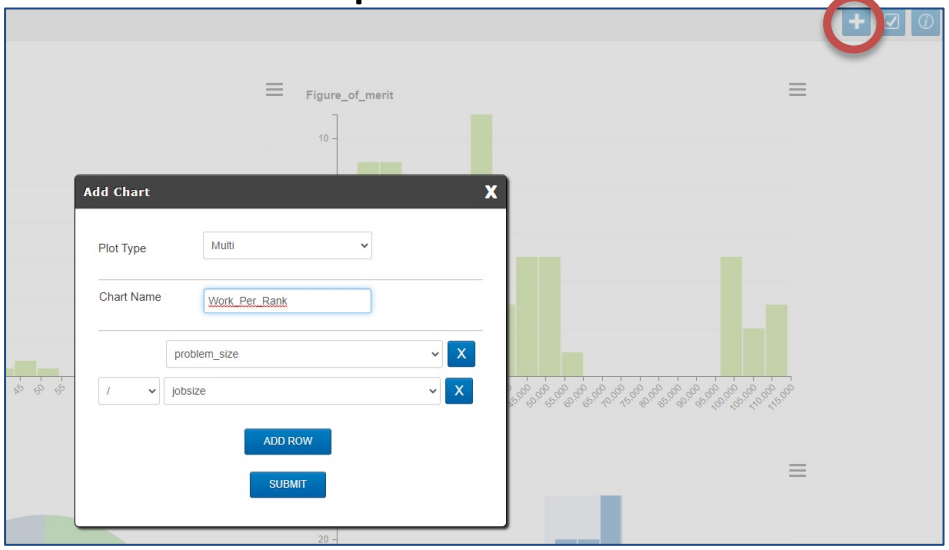

Performance Data can be Visualized Per-Run

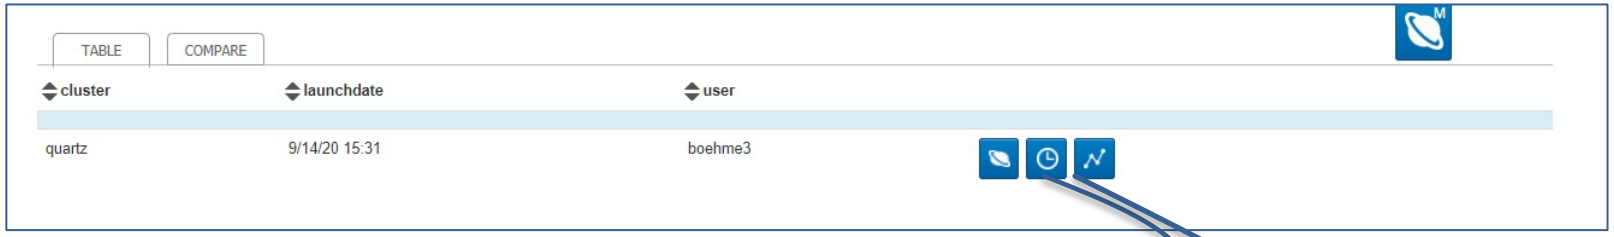

- FlameGraph: performance against code

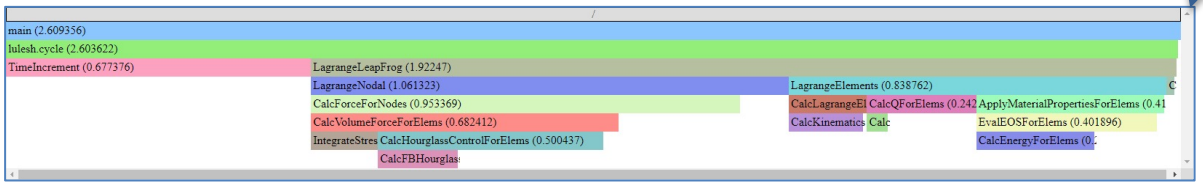

- Timeseries: performance against time

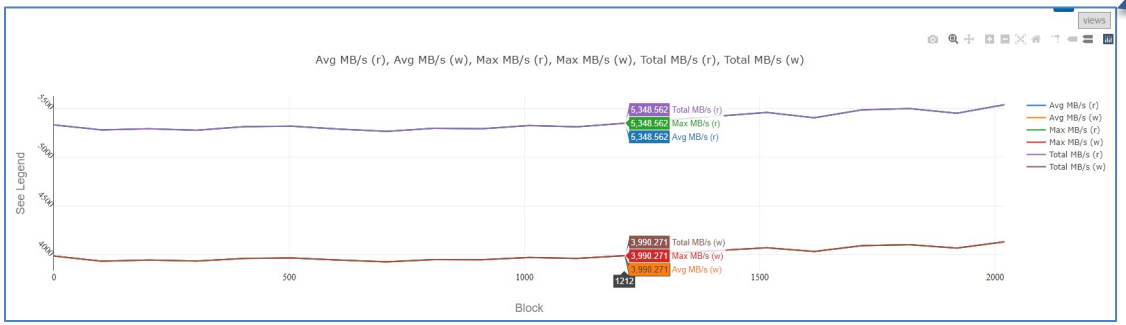

Performance Data can be Compared Across Runs

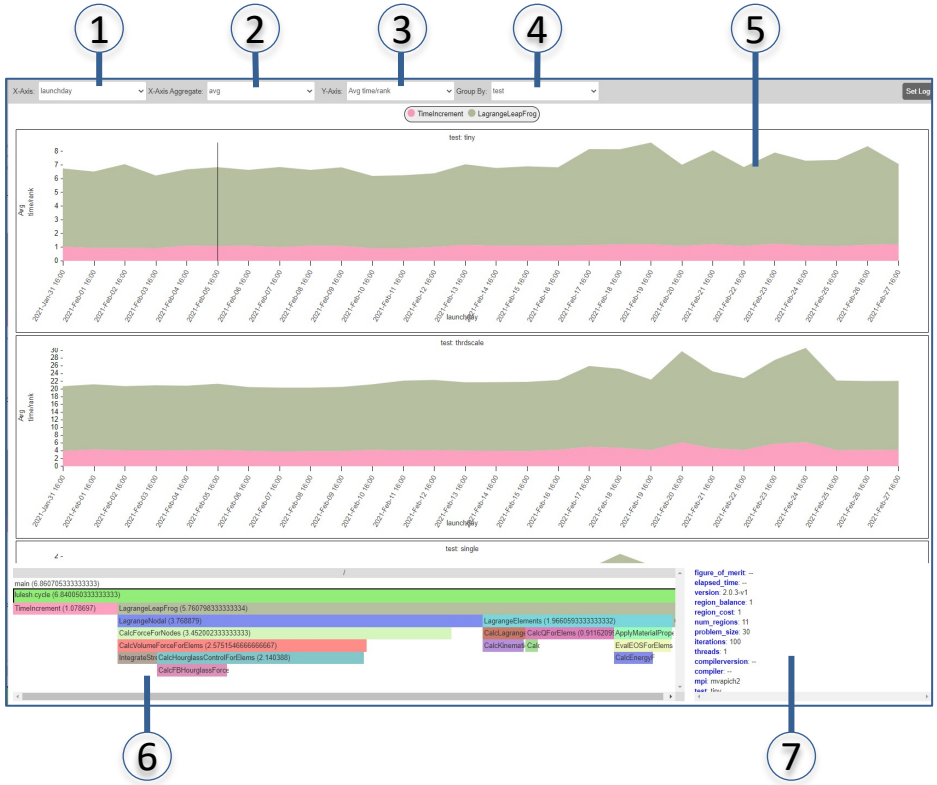

1. Metadata value on x-axis
2. Aggregation (min/max/avg…) if multiple runs at same point.
3. Metric to graph on y-axis.
4. Grouping metadata. Each unique value is a graph.
5. Stacked performance graph. Click to drill-down.
6. Flame graph for runs currently under cursor.
7. Metadata for runs currently under cursor.

Live Demo