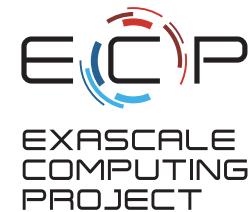


# Automated Application Performance Analysis with Caliper, SPOT, and Hatchet

2022 ECP Annual Meeting Tutorial

May 2, 2022



David Boehme, Stephanie Brink, Olga Pearce



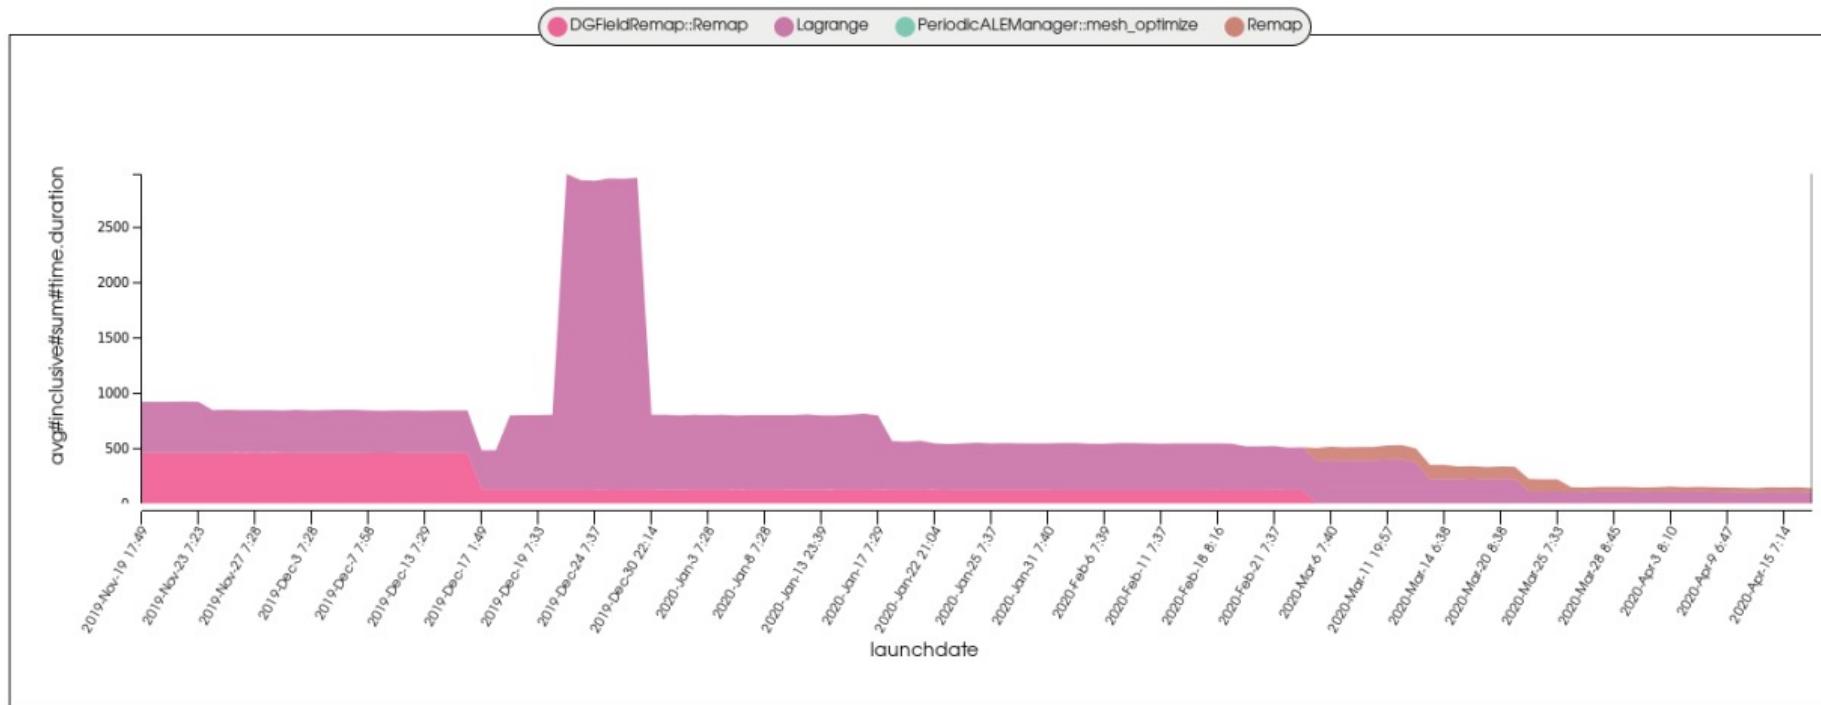
LLNL-PRES-821032

This work was performed under the auspices of the U.S. Department of Energy by Lawrence Livermore National Laboratory under contract DE-AC52-07NA27344. Lawrence Livermore National Security, LLC

Lawrence Livermore  
National Laboratory

# Building Automated Performance Analysis Workflows

Enabling performance analysis as a routine activity  
for HPC software development



Nightly test performance of a large physics code over 5 months

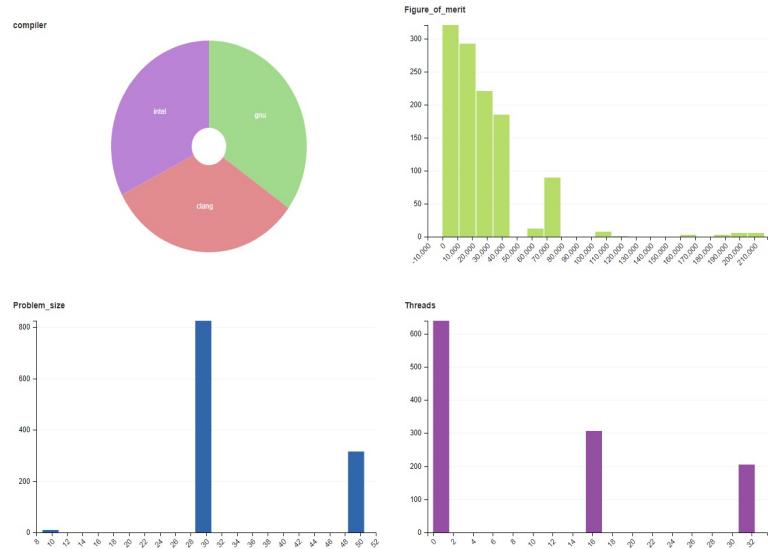
# Building Automated Performance Analysis Workflows



```
#include <caliper/cali.h>

void LagrangeElements(Domain& domain, Index_t numElem)
{
    CALI_CXX_MARK_FUNCTION;
// ...
```

Caliper:  
Instrumentation and Profiling



SPOT: Analysis of  
large collections of runs



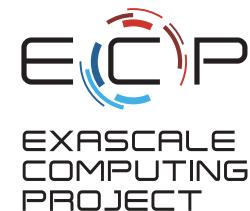
```
0.000 foo
|- 5.000 bar
  |- 5.000 baz
  |  |- 10.000 grault
|- 0.000 qux
  |- 5.000 quux
    |- 10.000 corge
      |- 5.000 bar
    |- 5.000 baz
    |  |- 10.000 grault
    |- 10.000 grault
    |- 15.000 garply
0.000 waldo
```

Hatchet:  
Call graph analysis in Python

# Caliper: A Performance Profiling Library

2022 ECP Annual Meeting: Tutorial

May 2, 2022



EXASCALE  
COMPUTING  
PROJECT



David Boehme  
Computer Scientist



LLNL-PRES-821032

This work was performed under the auspices of the U.S. Department of Energy by Lawrence Livermore National Laboratory under contract DE-AC52-07NA27344. Lawrence Livermore National Security, LLC

Lawrence Livermore  
National Laboratory

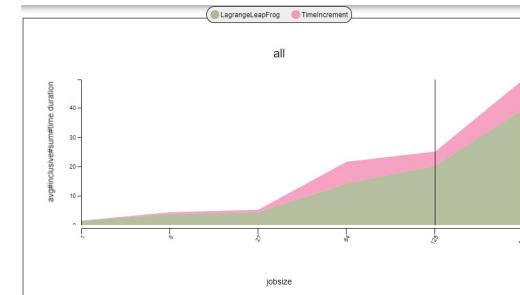
# Caliper: A Performance Profiling Library

- Integrates a performance profiler into your program
  - Profiling is always available
  - Simplifies performance profiling for application end users
- Common instrumentation interface
  - Provides program context information for other tools
- Advanced profiling features
  - MPI, CUDA, ROCm, Kokkos support; call-stack sampling; hardware counters; memory profiling

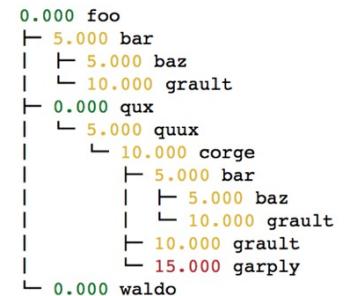
# Caliper Use Cases

- Lightweight always-on profiling
  - Performance summary report for each run
- Performance debugging
- Performance introspection
- Comparison studies across runs
  - Performance regression testing
  - Configuration and scaling studies
- Automated workflows

Performance reports					
Path	Min time/rank	Max time/rank	Avg time/rank	Time %	
main	0.000119	0.000119	0.000119	7.079120	
mainloop	0.000067	0.000067	0.000067	3.985723	
foo	0.000646	0.000646	0.000646	38.429506	
init	0.000017	0.000017	0.000017	1.011303	



Comparing runs



Debugging

# Performance Analysis with Caliper, SPOT and Hatchet

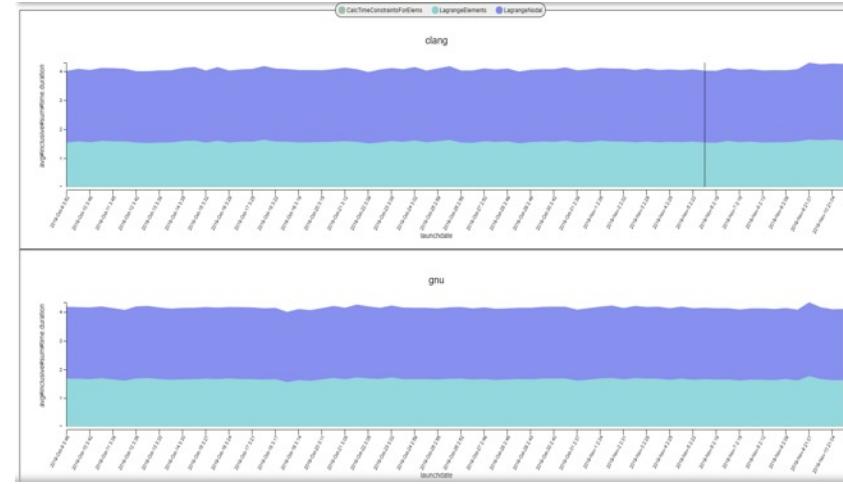


“spot” config

```
#include <caliper/caliper.h>

void LagrangeElements(Domain& domain,
Index_t numElem)
{
    CALI_CXX_MARK_FUNCTION;
// ...
```

Caliper:  
Instrumentation and Profiling



SPOT web frontend:  
Analysis of  
large collections of runs

hatchet-region-profile,  
hatchet-sample-profile

Pre-populated Jupyter  
notebooks



```
0.000 foo
|- 5.000 bar
  |- 5.000 baz
  |- 10.000 grault
|- 0.000 qux
  |- 5.000 quux
    |- 10.000 corge
      |- 5.000 bar
        |- 5.000 baz
        |- 10.000 grault
      |- 10.000 grault
      |- 15.000 garply
0.000 waldo
```

Hatchet:  
Call graph analysis in Python

# Materials, Contact & Links

- Tutorial materials: <https://github.com/daboehm/caliper-tutorial>

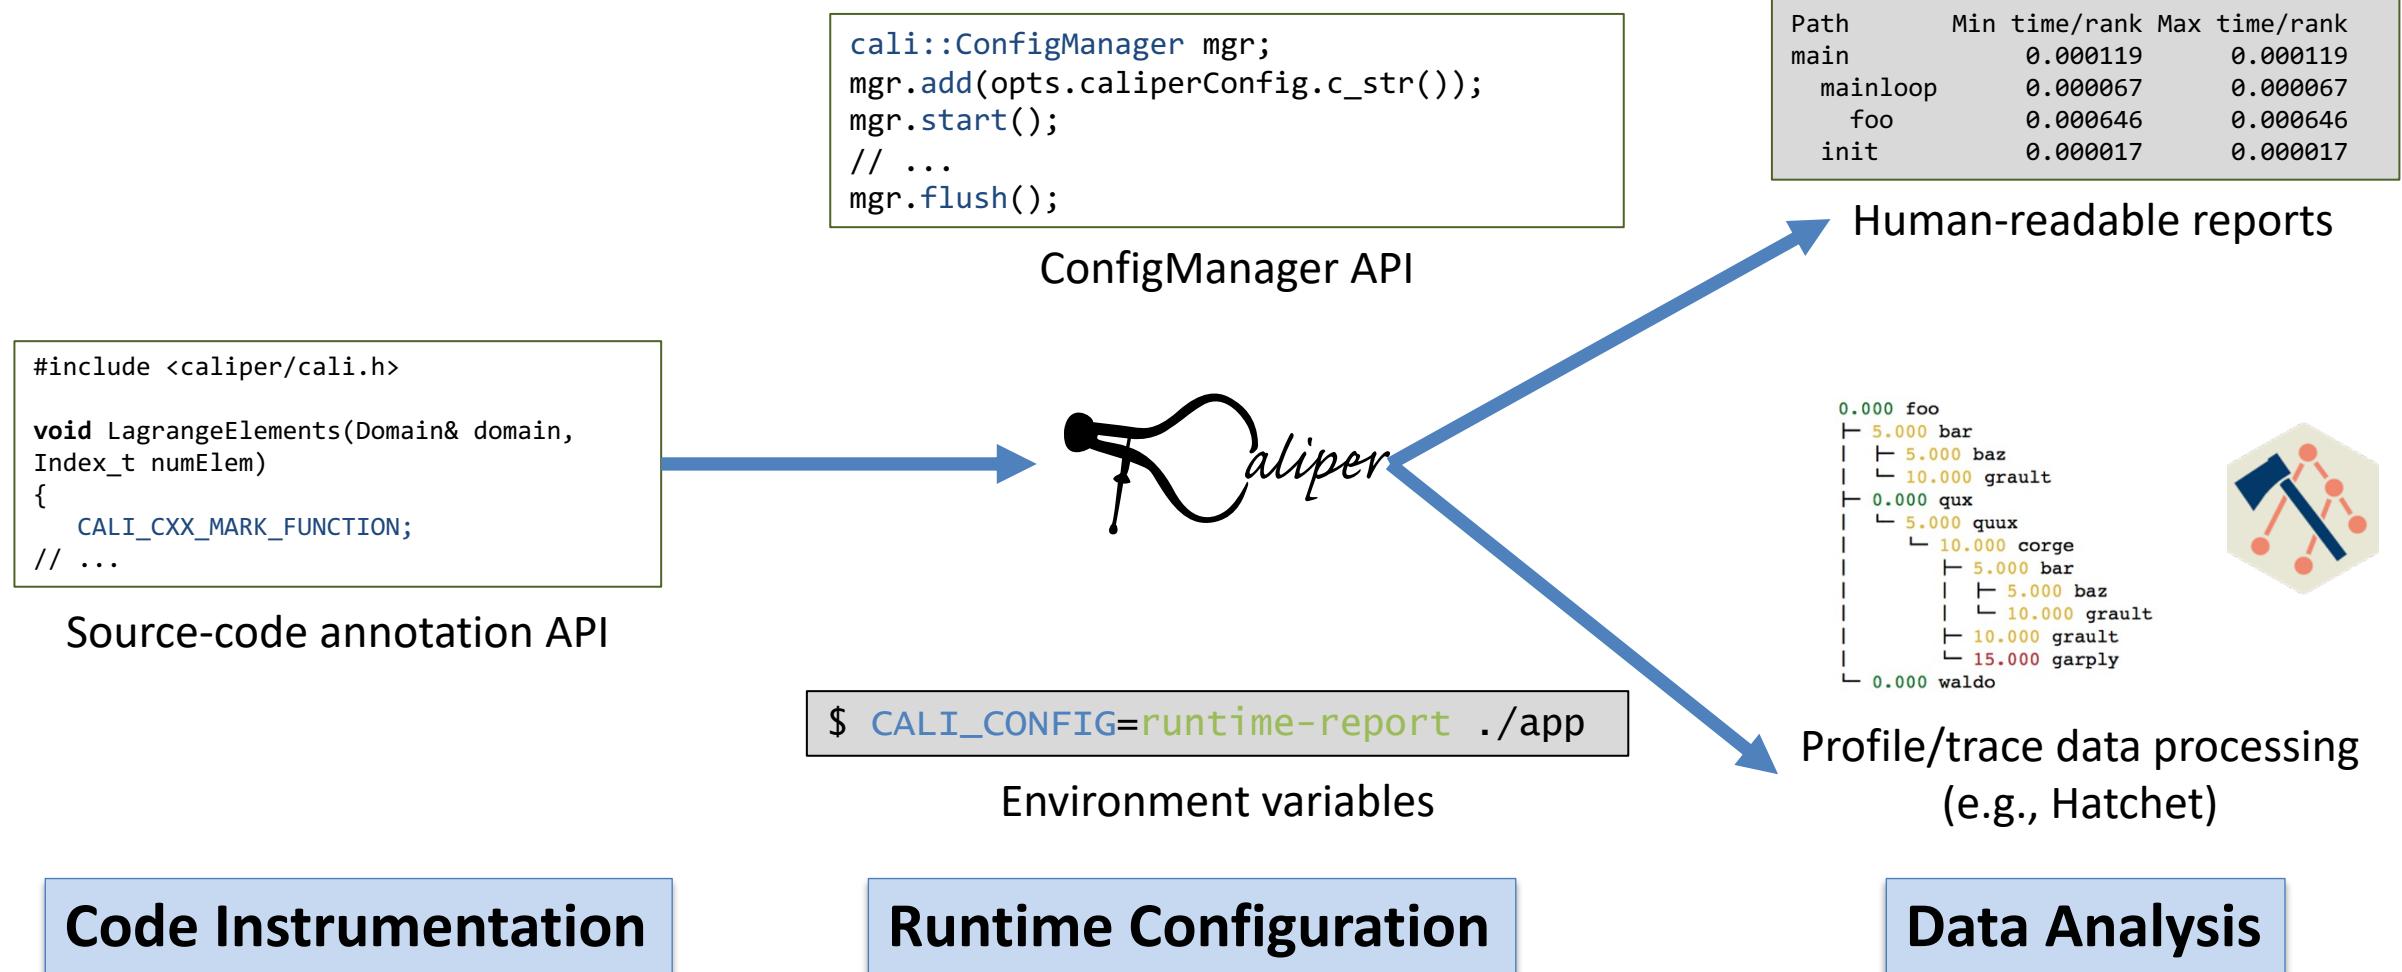
```
$ git clone --recursive https://github.com/daboehm/caliper-tutorial  
$ . setup-env.sh
```

- GitHub repository: <https://github.com/LLNL/Caliper>
- Documentation: <https://llnl.github.io/Caliper>
- GitHub Discussions: <https://github.com/LLNL/Caliper/discussions>
- Contact: David Boehme (boehme3@llnl.gov)

---

# Using Caliper

# Using Caliper: Workflow



# Region Profiling: Marking Code Regions

C/C++

```
#include <caliper/cali.h>

void main() {
    CALI_MARK_BEGIN("init");

    do_init();

    CALI_MARK_END("init");
}
```

Fortran

```
USE caliper_mod

CALL cali_begin_region('init')

CALL do_init()

CALL cali_end_region('init')
```

- Use annotation macros (C/C++) or functions to mark and name code regions

# Region Profiling: Best Practices

- Be selective: Instrument high-level program subdivisions (kernels, phases, ...)
- Be clear: Choose meaningful names
- Start small: Add instrumentation incrementally

```
RAJA::ReduceSum<RAJA::omp_reduce, double> ompdot(0.0);

CALI_MARK_BEGIN("dotproduct");

RAJA::forall<RAJA::omp_parallel_for_exec>(RAJA::RangeSegment(0, N), [=] (int i) {
    ompdot += a[i] * b[i];
});
dot = ompdot.get();

CALI_MARK_END("dotproduct");
```

Caliper annotations give meaningful names to high-level program constructs

# Region Profiling: Printing a Runtime Report

```
$ cd Caliper/build  
$ make cxx-example  
$ CALI_CONFIG=runtime-report ./examples/apps/cxx-example
```

Path	Min time/rank	Max time/rank	Avg time/rank	Time %
main	0.000119	0.000119	0.000119	7.079120
mainloop	0.000067	0.000067	0.000067	3.985723
foo	0.000646	0.000646	0.000646	38.429506
init	0.000017	0.000017	0.000017	1.011303

- Set the CALI\_CONFIG environment variable to access Caliper's built-in profiling configurations
- “runtime-report” measures, aggregates, and prints time in annotated code regions

# List of Caliper's Built-in Profiling Configurations

Config name	Description
runtime-report	Print a time profile for annotated regions
loop-report	Print summary and time-series information for loops
mpi-report	Print time spent in MPI functions
callpath-sample-report	Print time spent in functions using call-path sampling
event-trace	Record a trace of region enter/exit events in .cali format
hatchet-region-profile	Record a region time profile for processing with hatchet or cali-query
hatchet-sample-profile	Record a sampling profile for processing with hatchet or cali-query
spot	Record a time profile for the SPOT web visualization framework

Use `mpi-caliquery --help=configs` to list all built-in configs and their options

# Built-In Profiling Configurations: Configuration String Syntax

*Config name* specifies the kind of performance measurement

*Parameters* enable additional features, metrics, or output options

```
$ CALI_CONFIG="runtime-report(mem.highwatermark,output=stdout)" ./examples/apps/cxx-example
```

Path	Min time/rank	Max time/rank	Avg time/rank	Time %	Allocated MB
main	0.000179	0.000179	0.000179	2.054637	0.000047
mainloop	0.000082	0.000082	0.000082	0.941230	0.000016
foo	0.000778	0.000778	0.000778	8.930211	0.000016
init	0.000020	0.000020	0.000020	0.229568	0.000000

- Most Caliper measurement configurations have optional parameters to enable additional features or configure output settings

# Profiling Options: MPI Function Profiling

```
$ CALI_CONFIG=runtime-report,profile.mpi ./lulesh2.0
```

Path	Min time/rank	Max time/rank	Avg time/rank	Time %
MPI_Comm_dup	0.000034	0.003876	0.001999	0.10089
main	0.009013	0.010797	0.010173	0.51335
MPI_Reduce	0.000031	0.000049	0.000037	0.001886
lulesh.cycle	0.002031	0.002258	0.002085	0.105220
LagrangeLeapFrog	0.002158	0.002511	0.002227	0.112366
CalcTimeConstraintsForElms	0.015166	0.015443	0.015277	0.770922
CalcQForElms	0.058781	0.060196	0.059699	3.01254
CalcMonotonicQForElms	0.035331	0.041057	0.038496	1.942601
CommMonoQ	0.005280	0.006152	0.005544	0.279781
MPI_Wait	0.004182	0.084533	0.035324	1.78249
CommSend	0.006893	0.009062	0.008071	0.407298
MPI_Waitall	0.000986	0.001778	0.001343	0.067789
MPI_Isend	0.004564	0.005785	0.004930	0.248765
CommRecv	0.002265	0.002616	0.002341	0.118144
[...]				

The profile.mpi option measures time spent in MPI functions

# Profiling Options: CUDA Profiling

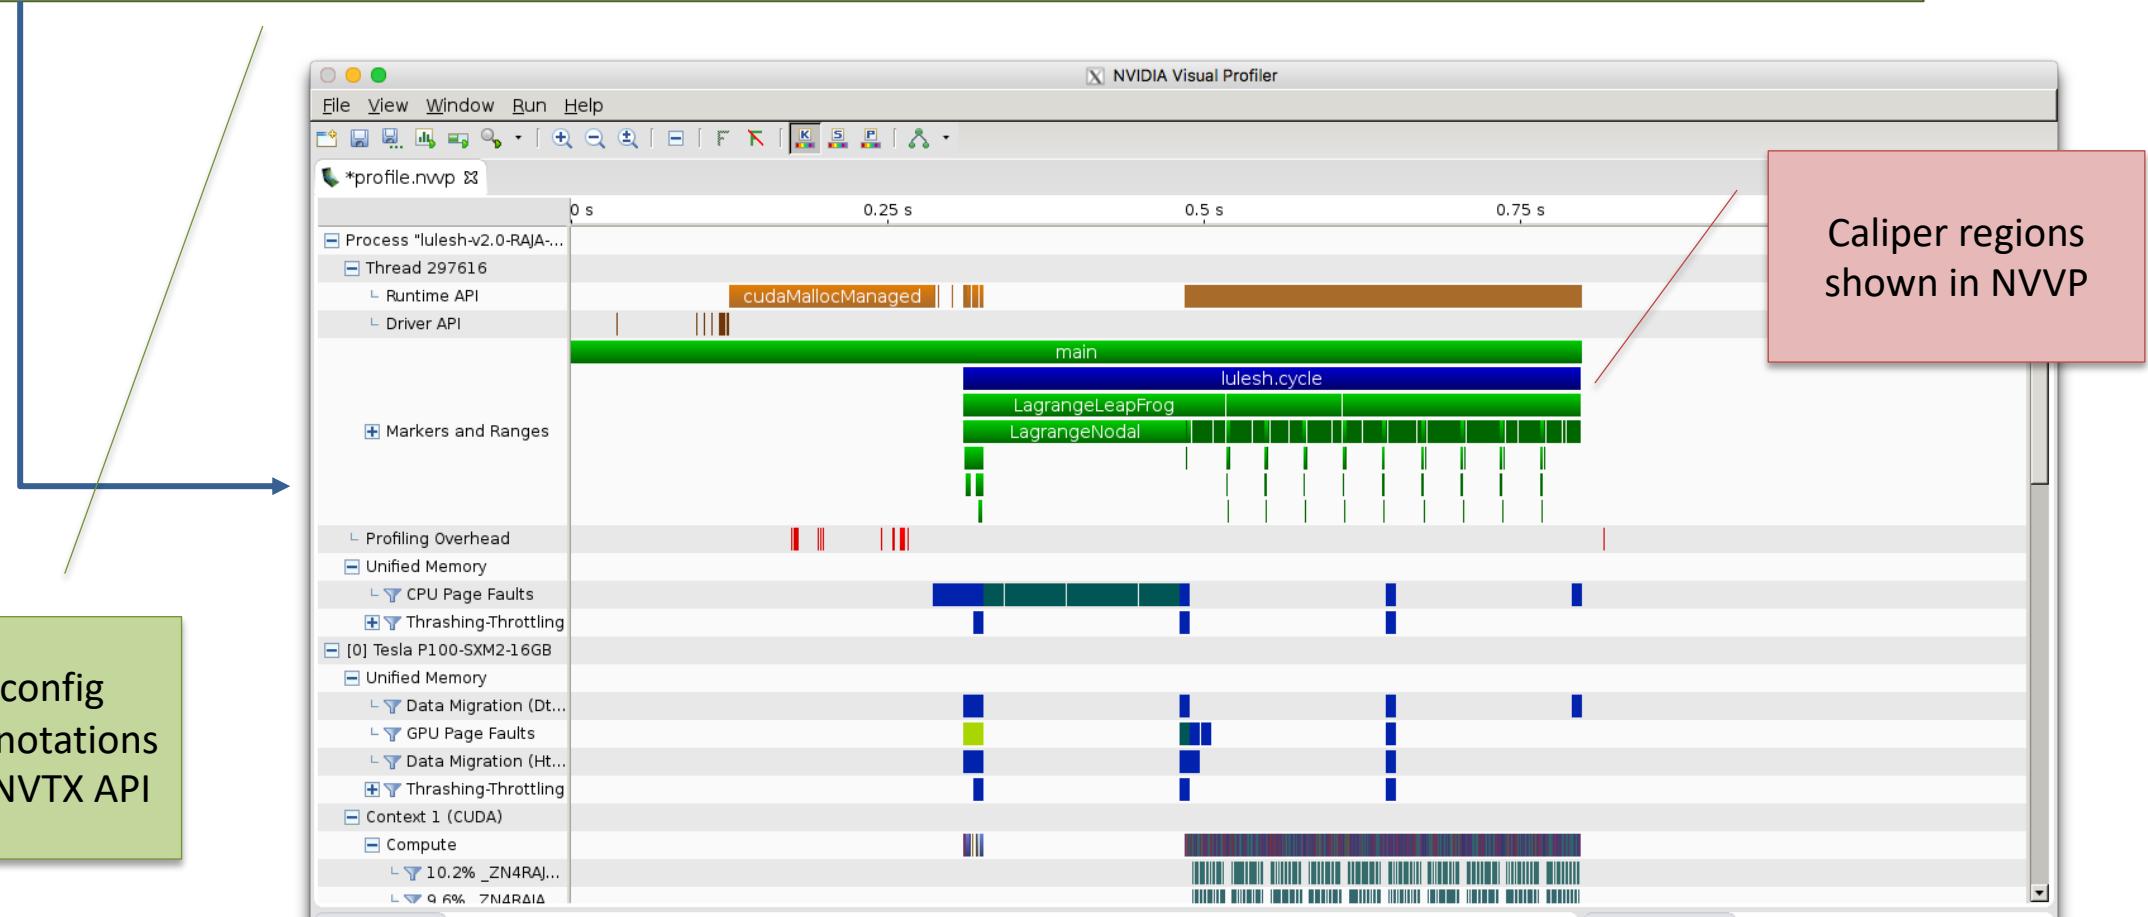
```
$ lrun -n 4 ./tea_leaf runtime-report,profile.cuda
```

Path	Min time/rank	Max time/rank	Avg time/rank	Time %
timestep_loop	0.000175	0.000791	0.000345	0.002076
[...]				
total_solve	0.000105	0.000689	0.000252	0.001516
solve	0.583837	0.617376	0.594771	3.581811
dot_product	0.000936	0.001015	0.000969	0.005837
cudaMalloc	0.000060	0.000066	0.000063	0.000382
internal_halo_update	0.077627	0.079476	0.078697	0.473925
halo_update	0.158597	0.161853	0.160023	0.963685
halo_exchange	1.502106	1.572522	1.532860	9.231136
cudaMemcpy	11.840890	11.871018	11.860343	71.424929
cudaLaunchKernel	1.177454	1.230816	1.211668	7.296865
cudaMemcpy	0.470123	0.471485	0.470596	2.834008
cudaLaunchKernel	0.658269	0.682566	0.673030	4.053100
[...]				

The profile.cuda option measures time in CUDA runtime API calls

# Forwarding Annotations to Third-Party Tools

```
$ CALI_CONFIG=nvtx nvprof <nvprof-opts> ./app
```



# Call Graph Analysis with the Hatchet Python Library

- Caliper records data for hatchet with hatchet-region-profile or hatchet-sample-profile

```
$ CALI_CONFIG=hatchet-sample-profile srun -n 8 ./lulesh2.0
```

```
>>> gf = hat.GraphFrame.from_caliper_json('/Users/boehme3/Documents/Data/lulesh_8x4_callpath-sample-profile.json')
>>> gf.subgraph_sum(['time'])
>>> gf = gf.filter(lambda x: x['name'] != '__restore_rt')
>>> gf = gf.filter(lambda x: x['name'].find('_omp_fn') == -1).squash()
>>> print(gf.tree())
      / \_ _ _ _ / \_ _ _ / \_ _ _ / \_ _ 
     /   \ \ / \ / \ / \ / \ / \ / \ / \ / \ / 
    / \ / / / / \ / / / / \ / / / / \ / / / 
   / \ / / \ \ / \ / \ / \ / \ / \ / \ / \ / \ / 
  / \ / / \ \ / \ / \ / \ / \ / \ / \ / \ / \ / \ / 
  v1.3.0

5.850 __clone
└ 5.850 start_thread
  └ 5.850 gomp_thread_start
    └ 0.070 CalcElemVolume(dou...t*, double const*)
      └ 0.005 UNKNOWN 4
        └ 0.075 cbrt
          └ 0.000 frexp
            └ 0.020 ldexp
              └ 0.010 scalbn
      └ 0.005 gomp_barrier_wait
      └ 2.545 gomp_barrier_wait_end
      └ 0.605 gomp_team_barrier_wait_end
```

Hatchet allows manipulation, computation, comparison, and visualization of call graph data

# Control Profiling Programmatically: The ConfigManager API

```
#include <caliper/cali.h>
#include <caliper/cali-manager.h>

int main(int argc, char* argv[])
{
    cali::ConfigManager mgr;
    mgr.add(argv[1]);
    if (mgr.error())
        std::cerr << mgr.error_msg() << "\n";

    mgr.start();
    // ...
    mgr.flush();
}
```

- Use ConfigManager to access Caliper's built-in profiling configurations

```
$ ./examples/apps/cxx-example -P runtime-report
```

- Now we can use command-line arguments or other program inputs to enable profiling

# Manual Configuration Allows Custom Analyses

```
cali-query -q "select alloc.label#cuhti.fault.addr as Pool,  
cuhti.uvm.kind as UVM\ Event,  
scale(cuhti.uvm.bytes,1e-6) as MB,  
scale(cuhti.activity.duration,1e-9) as Time  
group by  
prop:nested,alloc.label#cuhti.fault.addr,cuhti.uvm.kind  
where cuhti.uvm.kind format tree" trace.cali
```

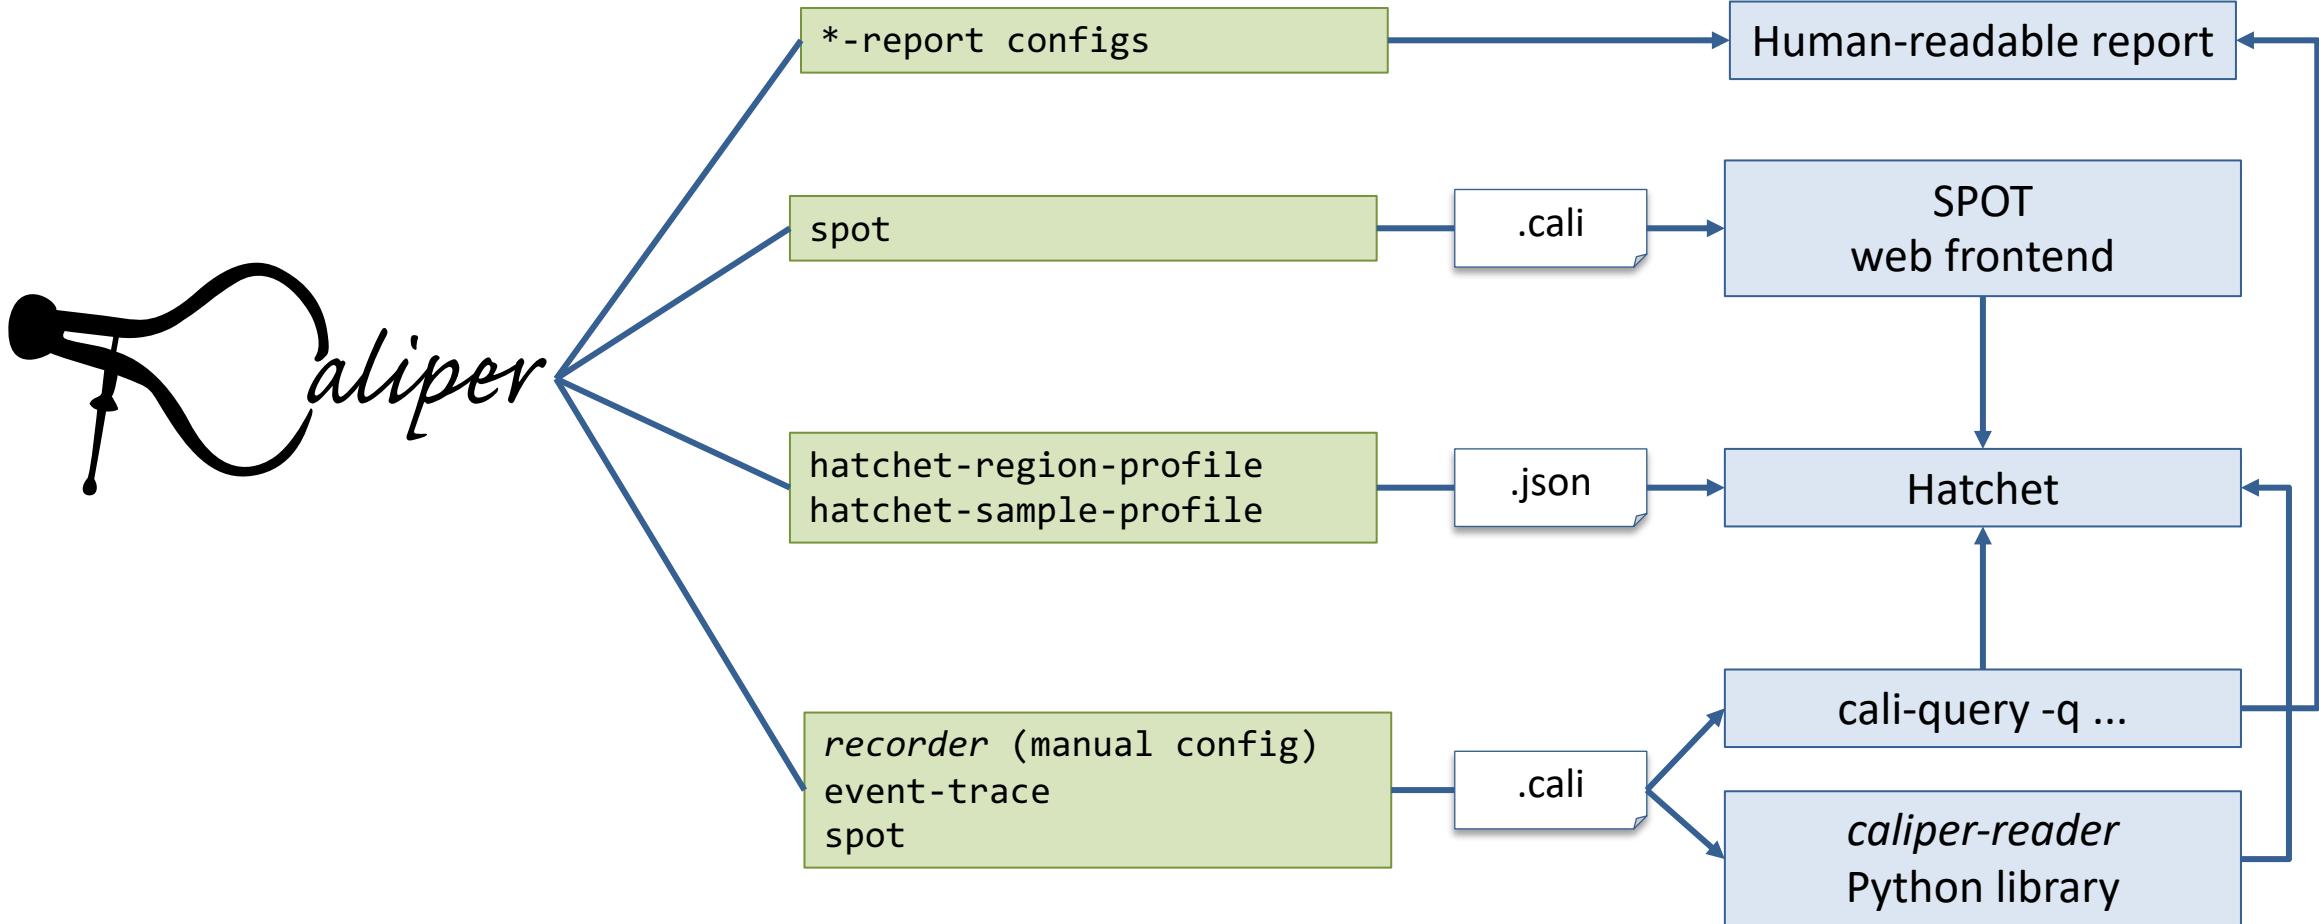
caliper.config

```
CALI_SERVICES_ENABLE=alloc,cuhti,cuhtitrace,mpi,trace,recorder  
CALI_ALLOC_RESOLVE_ADDRESSES=true  
CALI_CUHTI_CALLBACK_DOMAINS=sync  
CALI_CUHTITRACE_ACTIVITIES=uvm  
CALI_CUHTITRACE_CORRELATE_CONTEXT=false  
CALI_CUHTITRACE_FLUSH_ON_SNAPSHOT=true
```

Path				
main				
solve				
TIME_STEPPING				
enforceBC				
CURVI in EnforceBC				
CurviCartIC				
CurviCartIC::PART 3	Pool	UVM Event	MB	Time
curvilinear4sgwind	UM_pool	pagefaults.gpu		2.806946
curvilinear4sgwind	UM_pool	HtoD	7862.747136	0.232238
curvilinear4sgwind	UM_pool_temps	pagefaults.gpu		0.130167
curvilinear4sgwind	UM_pool	DtOH	9986.441216	0.378583
curvilinear4sgwind	UM_pool	pagefaults.cpu		

- Mapping CPU/GPU unified memory transfer events to Umpire memory pools in SW4

# Caliper Output Formats and Processing Workflows



---

# Recording Data for SPOT

# Recording Data for SPOT with Caliper and Adiak

```
#include <caliper/cali.h>

void LagrangeElements(Domain& domain,
Index_t numElem)
{
    CALI_CXX_MARK_FUNCTION;
// ...
```

Region  
instrumentation

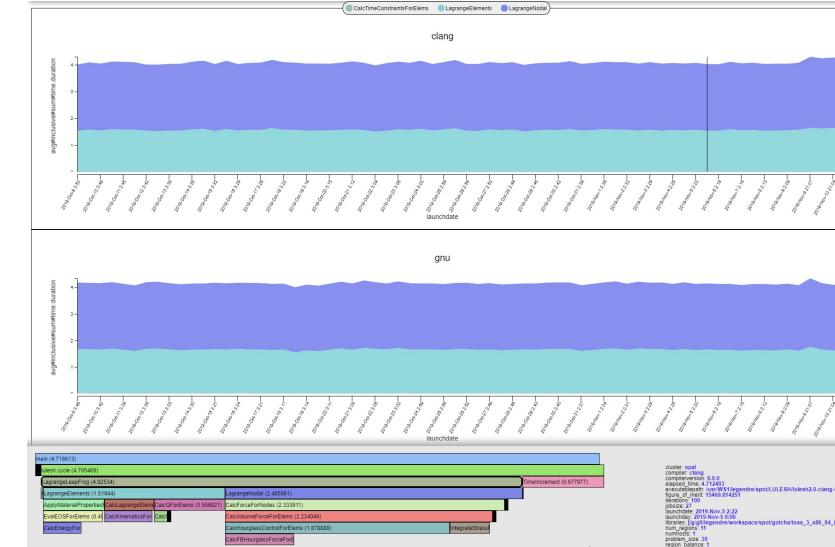
```
adiak::clusternode();
adiak::jobsizes();

adiak::value("iterations", opts.its);
adiak::value("problem_size", opts.nx);
adiak::value("num_regions", opts.numReg);
```

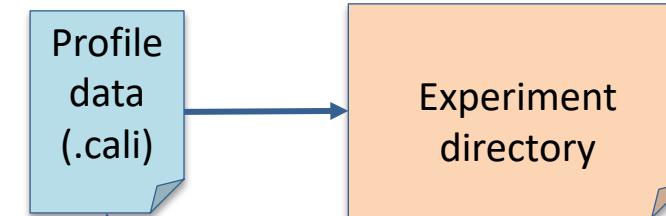
Metadata  
collection  
[adiak]

```
cali::ConfigManager mgr;
mgr.add(opts.caliperConfig.c_str());
mgr.start();
// ...
mgr.flush();
```

Caliper  
configuration



SPOT  
Web  
GUI



Run program with the “spot” profiling config

# Recording Program Metadata with the Adiak Library

TeaLeaf\_CUDA example [C++]

```
#include <adiak.hpp>

adiak::user();
adiak::launchdate();
adiak::jobsizes();

adiak::value("end_step", readInt(input, "end_step"));
adiak::value("halo_depth", readInt(input, "halo_depth"));

if (tl_use_ppcg) {
    adiak::value("solver", "PPCG");
// [...]
```

Use built-in Adiak functions to collect common metadata

Use key:value functions to collect program-specific data

- Use the [Adiak](#) C/C++ library to record program metadata
  - Environment info (user, launchdate, system name, ...)
  - Program configuration (input problem description, problem size, ...)
- Enables performance comparisons across runs. Required for SPOT.

# Adiak: Built-in Functions for Common Metadata

```
adiak_user();          /* user name */  
adiak_uid();          /* user id */  
adiak_launchdate();   /* program start time (UNIX timestamp) */  
adiak_executable();   /* executable name */  
adiak_executablepath();/* full executable file path */  
adiak_cmdline();      /* command line parameters */  
adiak_hostname();     /* current host name */  
adiak_clustername();  /* cluster name */  
  
adiak_job_size();     /* MPI job size */  
adiak_hostlist();     /* all host names in this MPI job */  
  
adiak_walltime();     /* wall-clock job runtime */  
adiak_cputime();      /* job cpu runtime */  
adiak_systime();      /* job sys runtime */
```

- Adiak comes with built-in functions to collect common environment metadata
- SPOT requires at least `launchdate`

# Adiak: Recording Custom Key-Value Data in C++

C++

```
#include <adiak.hpp>

vector<int> ints { 1, 2, 3, 4 };
adiak::value("myvec", ints);

adiak::value("myint", 42);
adiak::value("mydouble", 3.14);
adiak::value("mystring", "hi");

adiak::value("mypath", adiak::path("/dev/null"));
adiak::value("compiler", adiak::version("gcc@8.3.0"));
```

- Adiak supports many basic and structured data types
  - Strings, integers, floating point, lists, tuples, sets, ...
- `adiak::value()` records key:value pairs with overloads for many data types

# Adiak: Recording Custom Key-Value Data in C

C

```
#include <adiak.h>

int ints[] = { 1, 2, 3, 4 };
adiak_nameval("myvec",     adiak_general, NULL, "[%d]", ints, 4);

adiak_nameval("myint",     adiak_general, NULL, "%d", 42);
adiak_nameval("mydouble",   adiak_general, NULL, "%f", 3.14);
adiak_nameval("mystring",   adiak_general, NULL, "%s", "hi");

adiak_nameval("mypath",    adiak_general, NULL, "%p", "/dev/null");
adiak_nameval("compiler",   adiak_general, NULL, "%v", "gcc@8.3.0");
```

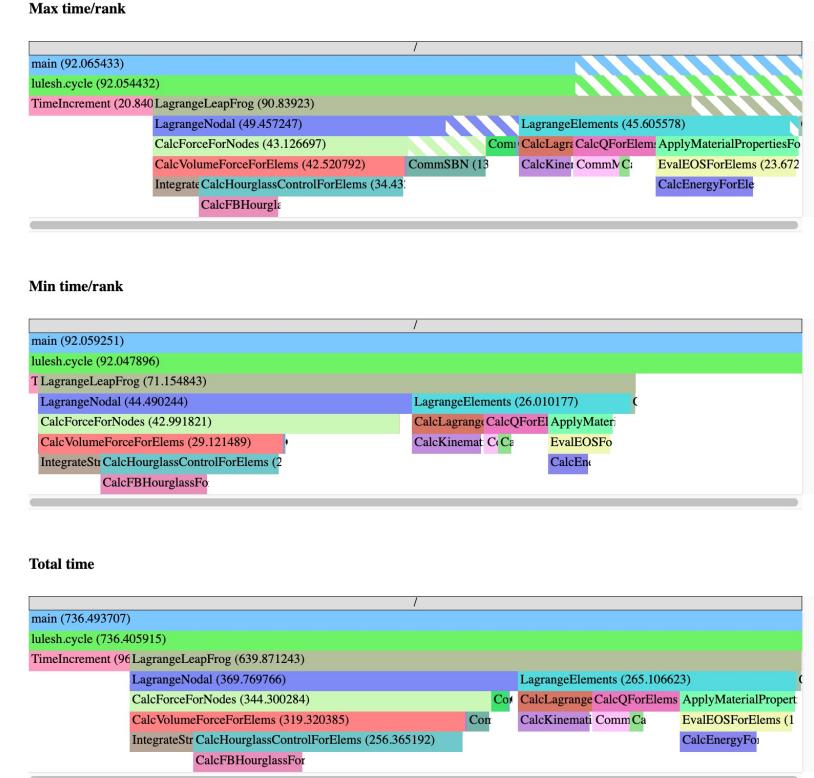
- In C, `adiak_nameval()` uses `printf()`-style descriptors to determine data types

# The spot config: Region Profiling

```
$ CALI_CONFIG=spot,profile.mpi ./lulesh2.0
```

```
$ ls *.cali  
210304-17175150010.cali
```

- “spot” records and aggregates time spent in instrumented regions, like runtime-report
- Supports many profiling options (e.g., MPI function profiling)
- Collect profiling output (.cali files) in a directory for analysis in SPOT



SPOT region profile flame graphs

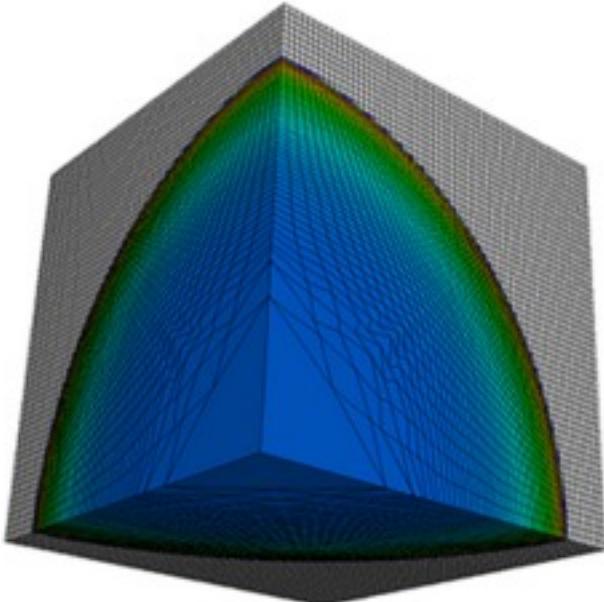
---

# Example: Caliper and Adiak in LULESH

# Modified LULESH Proxy App with Caliper and Adiak Support

<https://github.com/daboehm/LULESH/tree/adiak-caliper-support>

```
$ mpirun -n 8 ./lulesh2.0 -P runtime-report,profile.mpi
```



Path	Min time/rank	Max time/rank	Avg time/rank	Time
%				
MPI_Comm_dup	0.000034	0.003876	0.001999	0.10089
main	0.009013	0.010797	0.010173	0.51335
MPI_Reduce	0.000031	0.000049	0.000037	0.001886
lulesh.cycle	0.002031	0.002258	0.002085	0.105220
LagrangeLeapFrog	0.002158	0.002511	0.002227	0.112366
CalcTimeConstraintsForElems	0.015166	0.015443	0.015277	0.770922
CalcQForElems	0.058781	0.060196	0.059699	3.01254
CalcMonotonicQForElems	0.035331	0.041057	0.038496	1.942601
CommMonoQ	0.005280	0.006152	0.005544	0.279781
MPI_Wait	0.004182	0.084533	0.035324	1.78249
CommSend	0.006893	0.009062	0.008071	0.407298
MPI_Waitall	0.000986	0.001778	0.001343	0.067789
MPI_Isend	0.004564	0.005785	0.004930	0.248765
CommRecv	0.002265	0.002616	0.002341	0.118144
[...]				

# LULESH Example: Region Annotations

```
void CalcLagrangeElements(Domain& domain)
{
    CALI_CXX_MARK_FUNCTION;
    ...
}
```

Function annotation in LULESH

- Top-level functions provide meaningful basis for performance analysis in LULESH
- Annotated 17 out of 39 computational functions and 5 communication functions

# LULESH Example: Main Loop Annotation

```
CALI_CXX_MARK_LOOP_BEGIN(cycleloop, "lulesh.cycle");

while((locDom->time() < locDom->stoptime()) && (locDom->cycle() < opts.its)) {
    CALI_CXX_MARK_LOOP_ITERATION(cycleloop, locDom->cycle());
    // ...
}

CALI_CXX_MARK_LOOP_END(cycleloop);
```

Main loop annotation in LULESH

- Annotation of the main time-stepping loop and iterations for loop profiling

# LULESH Example: Initialization and ConfigManager

```
adiak::init(adiak_comm_p);

cali::ConfigManager mgr;
if (!opts.caliperConfig.empty())
    mgr.add(opts.caliperConfig.c_str());

if (mgr.error())
    std::cerr << "Caliper config parse error: " << mgr.error_msg() << std::endl;

mgr.start();
// ...
mgr.flush();
MPI_Finalize();
```

ConfigManager setup in LULESH

- Profiling control via ConfigManager API
- Modified LULESH command-line parsing code to read Caliper config string (not shown)

# LULESH Example: Recording Metadata With Adiak

```
void RecordGlobals(const cmdLineOpts& opts, int num_threads)
{
    adiak::user();
    adiak::launchdate();
    adiak::executablepath();
    adiak::libraries();
    adiak::cmdline();
    adiak::clustername();
    adiak::jobsize();

    adiak::value("threads", num_threads);
    adiak::value("iterations", opts.its);
    adiak::value("problem_size", opts.nx);
    adiak::value("num_regions", opts.numReg);
    adiak::value("region_cost", opts.cost);
    adiak::value("region_balance", opts.balance);
}
```

Recording environment and LULESH config

```
void VerifyAndWriteFinalOutput(...)
{
    // ...
    adiak::value("elapsed_time", elapsed_time);
    adiak::value("figure_of_merit", 1000.0/grindTime2);
}
```

Recording global performance metrics at program end

- Adiak calls record environment info, LULESH configuration options, and global performance metrics

# LULESH Example: Build System Modifications

```
find_package(caliper REQUIRED)
find_package(adiak REQUIRED)

# ...

add_executable(${LULESH_EXEC} ${LULESH_SOURCES})

target_include_directories(${LULESH_EXEC} PRIVATE ${caliper_INCLUDE_DIR} ${adiak_INCLUDE_DIRS})
target_link_libraries(${LULESH_EXEC} caliper adiak)
```

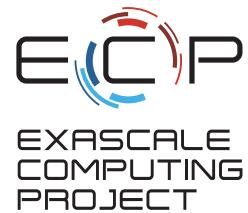
CMakeLists.txt

- Using caliper and adiak `find_package()` support in LULESH CMake script

# SPOT Tutorial Slides

2022 ECP Annual Meeting

May 2, 2022



Matthew LeGendre



LLNL-PRES-821032

This work was performed under the auspices of the U.S. Department of Energy by Lawrence Livermore National Laboratory under contract DE-AC52-07NA27344. Lawrence Livermore National Security, LLC

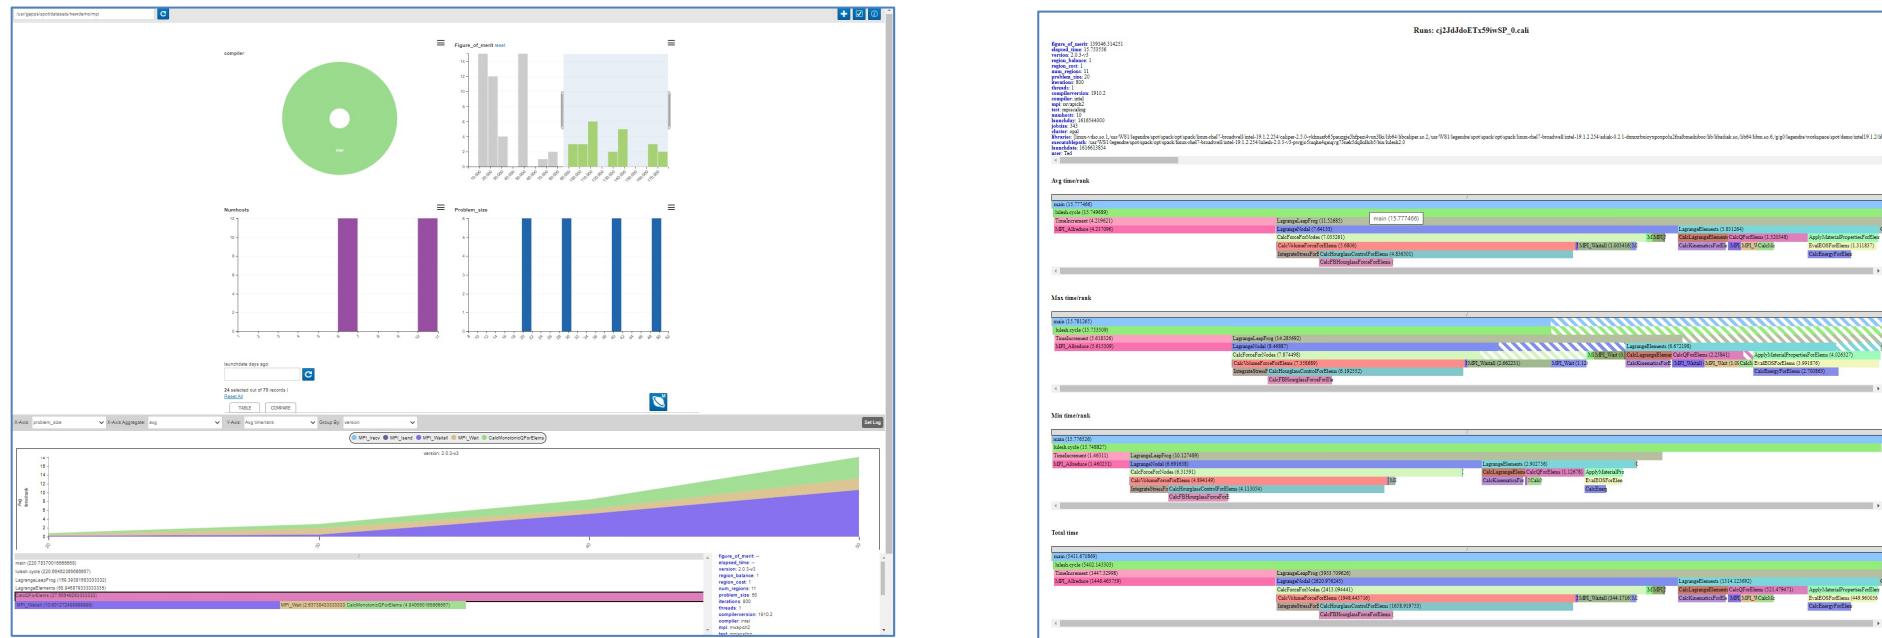
 Lawrence Livermore  
National Laboratory

# Getting SPOT

- LLNL installation: <https://lc.llnl.gov/spot2>
- Docker container: [https://github.com/LLNL/spot2\\_container](https://github.com/LLNL/spot2_container)

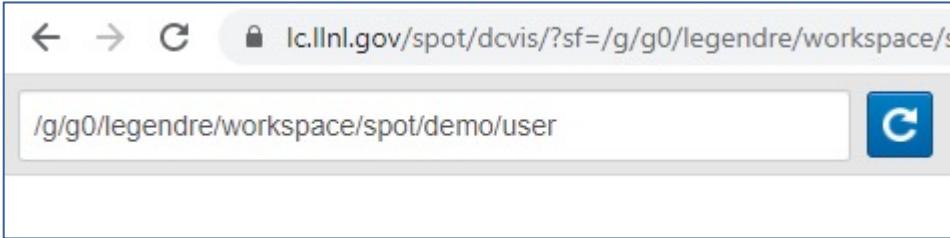
# SPOT: Software Performance and Analysis Tracking

- SPOT is visualization for **collections of performance profiles**. Common uses cases are:
  - Performance comparison of nightly tests. Look for performance regressions.
  - Performance tracking of developer changes. Run an MPI scaling study.
  - Collect performance profiles from users. Understand how users run and hit performance issues.



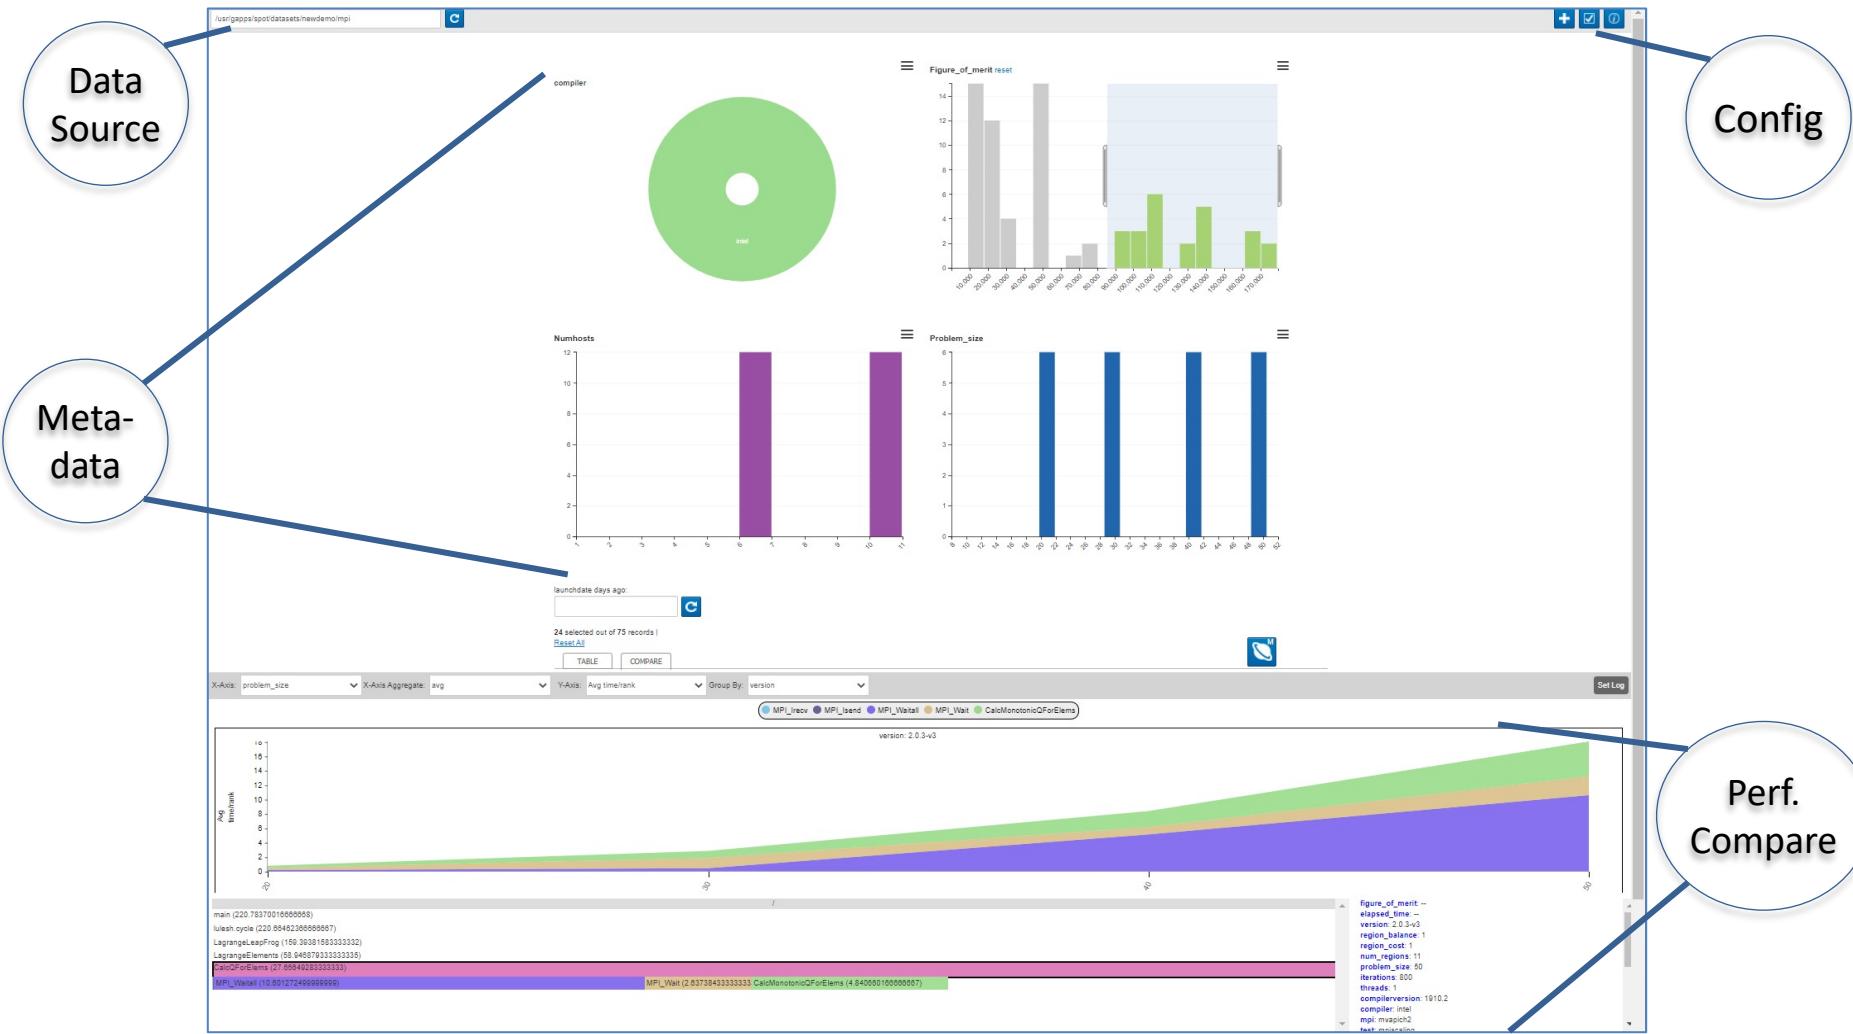
# SPOT's Workflow

1. Point SPOT webpage at directory containing \*.cali files.

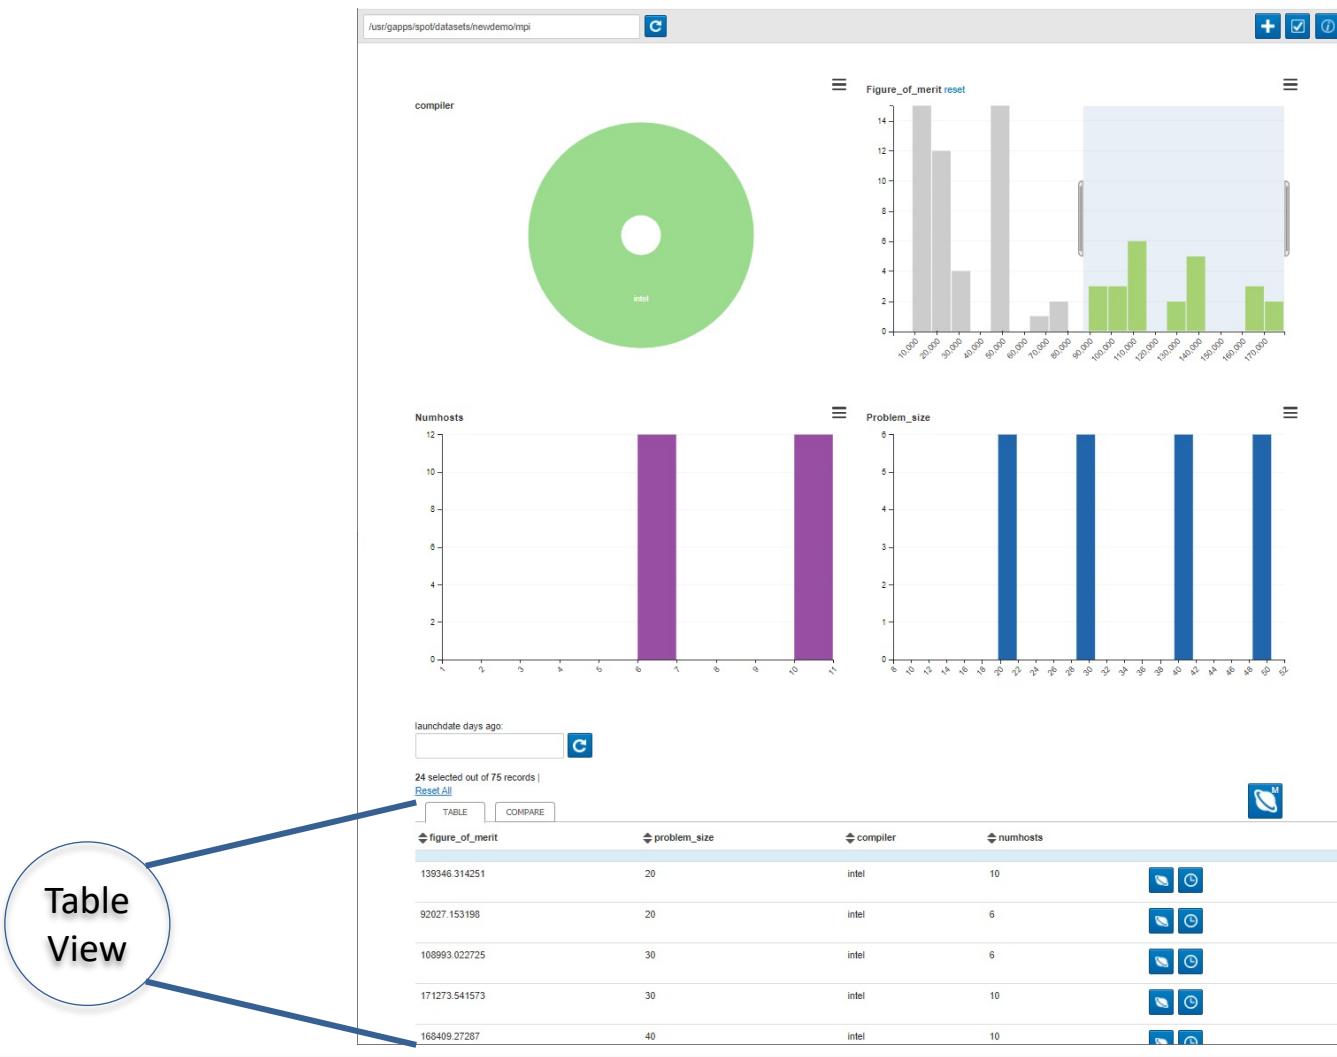


2. Select “interesting” sets of runs to analyze
  - Don’t compare performance of 1d test run. vs 3d multi-physics run.
3. Run common analysis with SPOT.
  - i.e., Graph history of performance changes
  - i.e., View performance context trees
4. Run advanced analysis with Hatchet.

# SPOT's Landing Page



# SPOT's Landing Page



# SPOT is Based on CrossFilter

TABLE		COMPARE								
◆ figure_of_merit	◆ elapsed_time	◆ version	◆ problem_size	◆ threads	◆ numhosts	◆ jobsize	◆ launchdate	◆ user		
9657.624636	2.60561	2.0.3-v1	10	1	2	27	2/3/21 8:11	bob		
24056.116646	39.672571	2.0.3-v1	40	1	2	64	2/20/21 8:11	alice		
10150.910684	10.306957	2.0.3-v3	50	1	1	27	2/18/21 8:11	alice		
24072.920076	24.671706	2.0.3-v1	40	1	2	64	2/14/21 8:11	alice		
23945.74189	36.081572	2.0.3-v3	50	1	2	64	3/4/21 8:11	alice		
23258.383892	39.211667	2.0.3-v2	50	1	2	64	2/16/21 8:11	alice		
15125.493765	10.264128	2.0.3-v3	50	1	2	27	2/19/21 8:11	alice		
23189.230912	25.529092	2.0.3-v3	50	1	2	64	2/22/21 8:11	alice		
24863.815433	35.418538	2.0.3-v2	40	1	2	64	2/12/21 8:11	alice		
23622.598644	33.527217	2.0.3-v3	50	1	2	64	2/24/21 8:11	alice		

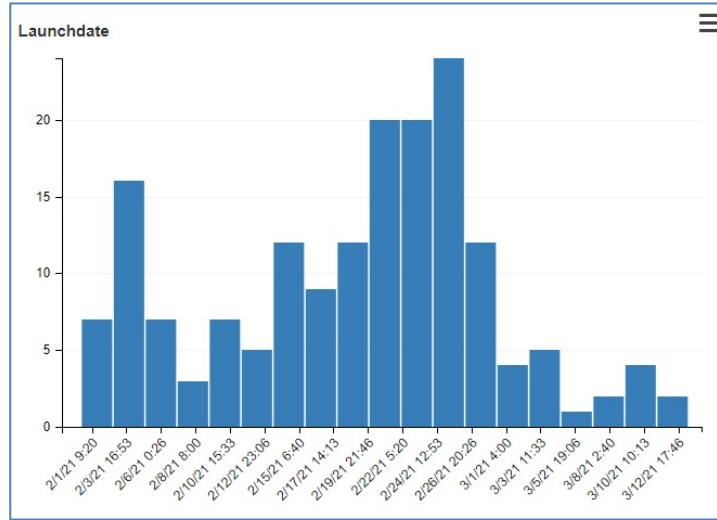
SPOT's Crossfilter: <https://dc-js.github.io/dc.js/>  
D3-Based Crossfilter: <https://square.github.io/crossfilter/>

- CrossFilter is good at finding correlations in data.
- Crossfilter helps you manage/slice your data to make interesting comparisons.

# Each App Run has Two Types of Data

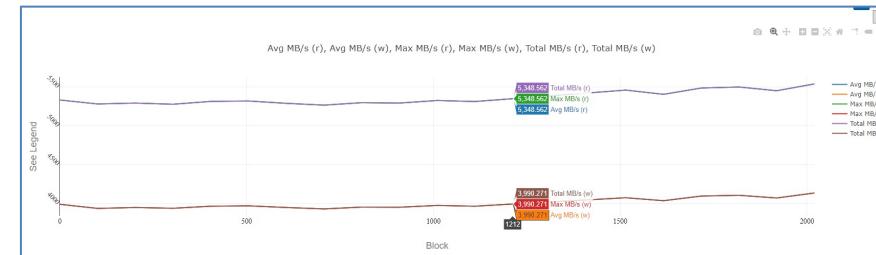
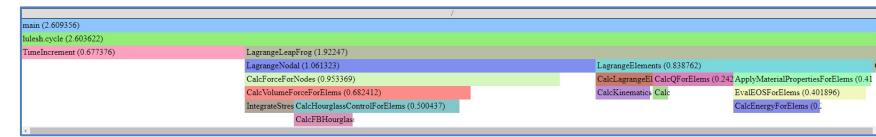
## ▪ Metadata

- Name/value data about job.
  - User, host, FOM, input parameters, ...
- Displayed as histograms



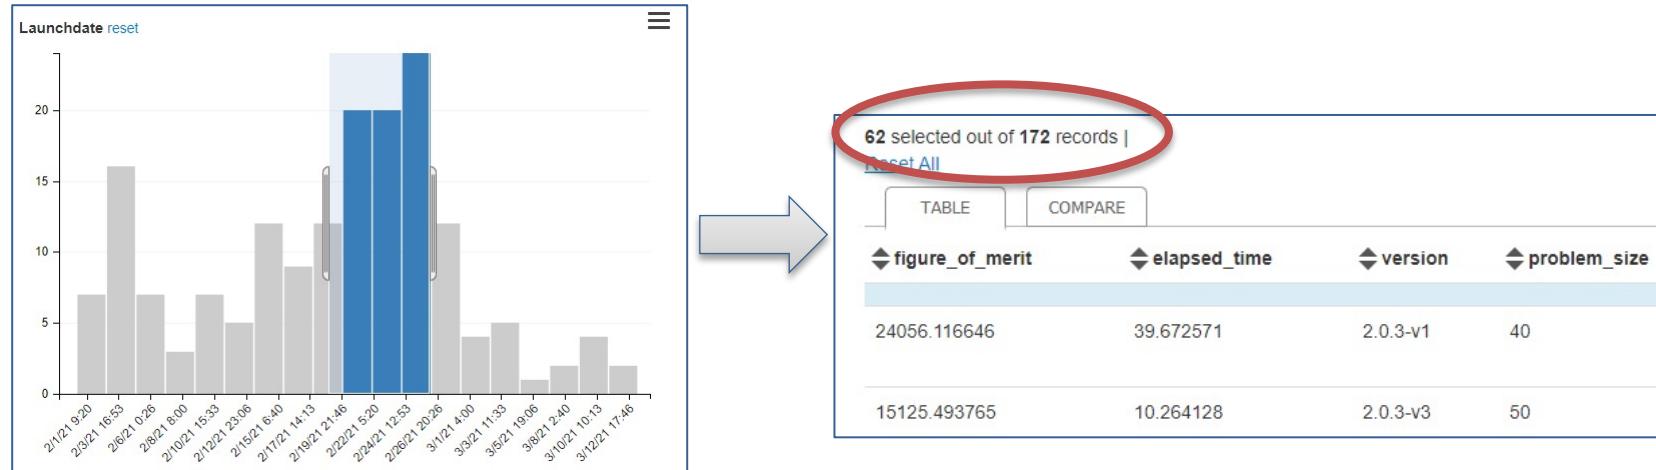
## ▪ Performance Data

- Can be different metrics.
  - Avg Time/Rank, Max Time/Rank, bandwidth, ...
- Displayed as a flame graph or timeseries.



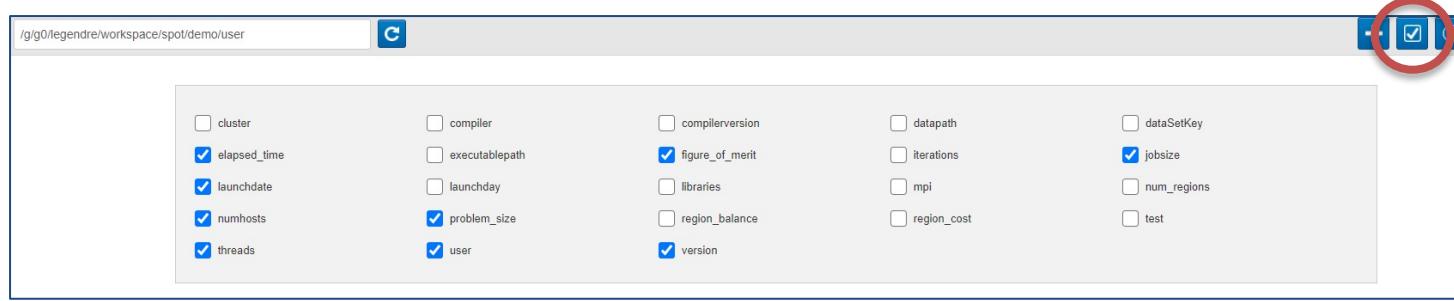
# Metadata Can be Filtered

- Select regions of metadata histograms to filter.
  - Filtering in one histogram updates all histograms range.
    - E.g., Filtering Launchdate to 'Feb 19-Feb26' will update the 'Users' histogram to only show users who ran between those dates.
  - While filtered, comparisons and Jupyter notebooks will operate on the filtered data.

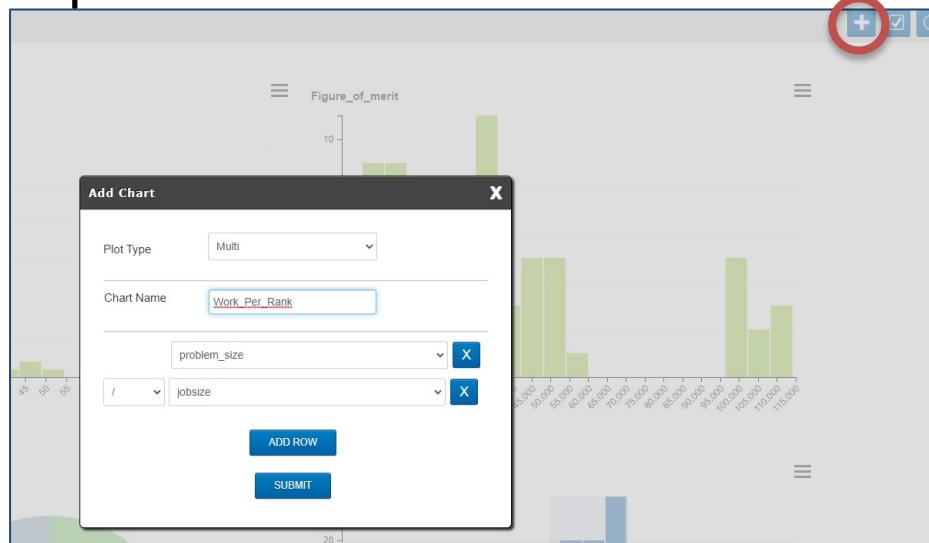


# Metadata is Configurable

- Can show/hide metadata:



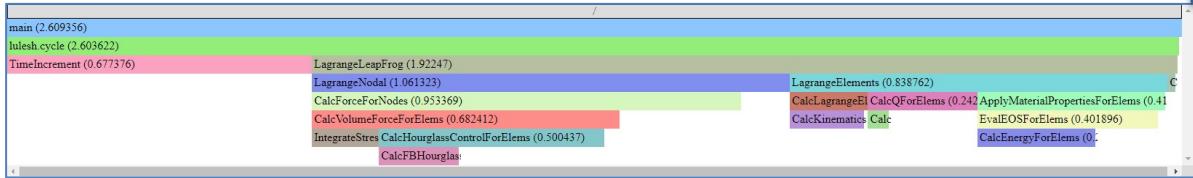
- Can create composite metadata:



# Performance Data can be Visualized Per-Run



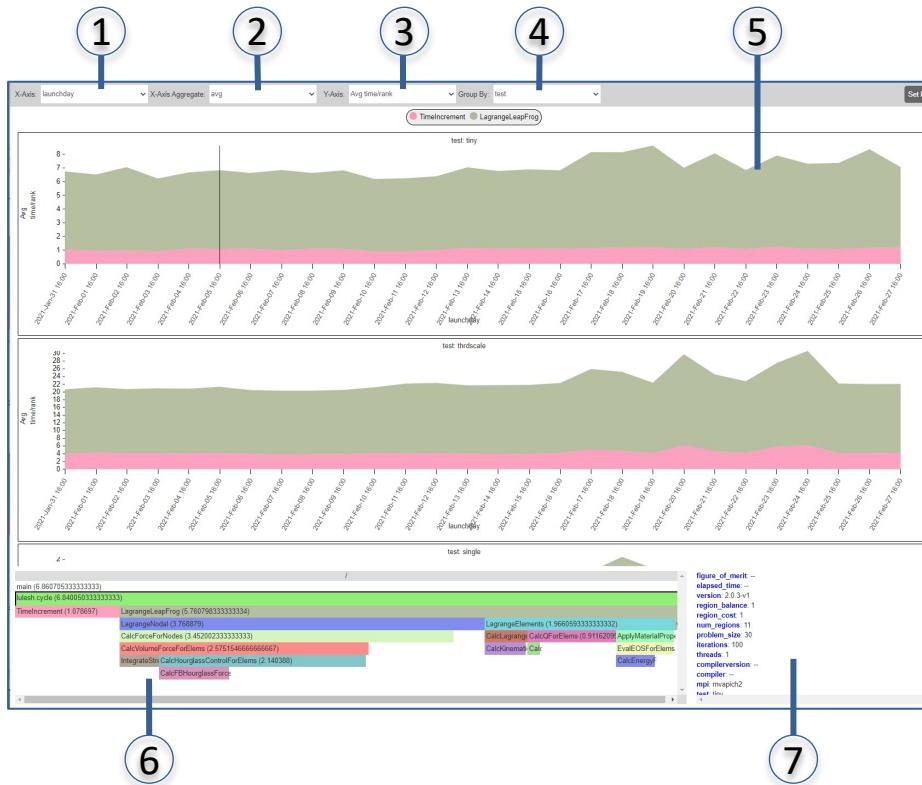
- FlameGraph: performance against code



- Timeseries: performance against time



# Performance Data can be Compared Across Runs

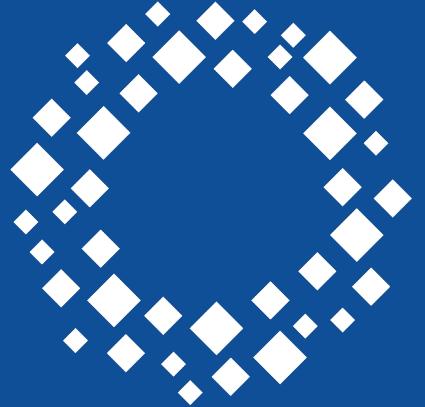


1. Metadata value on x-axis
2. Aggregation (min/max/avg...) if multiple runs at same point.
3. Metric to graph on y-axis.
4. Grouping metadata. Each unique value is a graph.
5. Stacked performance graph. Click to drill-down.
6. Flame graph for runs currently under cursor.
7. Metadata for runs currently under cursor.

# Live Demo



**Lawrence Livermore  
National Laboratory**



**CASC**

Center for Applied  
Scientific Computing

This research was supported by the Exascale Computing Project (17-SC-20-SC), a joint project of the U.S. Department of Energy's Office of Science and National Nuclear Security Administration, responsible for delivering a capable exascale ecosystem, including software, applications, and hardware technology, to support the nation's exascale computing imperative.

**Disclaimer**

This document was prepared as an account of work sponsored by an agency of the United States government. Neither the United States government nor Lawrence Livermore National Security, LLC, nor any of their employees makes any warranty, expressed or implied, or assumes any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represents that its use would not infringe privately owned rights. Reference herein to any specific commercial product, process, or service by trade name, trademark, manufacturer, or otherwise does not necessarily constitute or imply its endorsement, recommendation, or favoring by the United States government or Lawrence Livermore National Security, LLC. The views and opinions of authors expressed herein do not necessarily state or reflect those of the United States government or Lawrence Livermore National Security, LLC, and shall not be used for advertising or product endorsement purposes.

