



# Performance Analysis with Hatchet

2 May 2022

Stephanie Brink, Olga Pearce



LLNL-PRES-813307

This work was performed under the auspices of the U.S. Department of Energy by Lawrence Livermore National Laboratory under contract DE-AC52-07NA27344. Lawrence Livermore National Security, LLC

 Lawrence Livermore  
National Laboratory

# Getting Hatchet Tutorial Materials

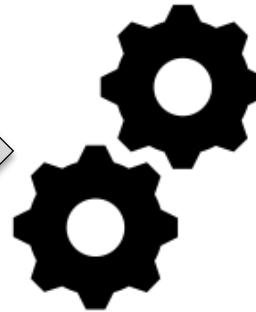
- The SPOT container includes a sample Jupyter notebook, Hatchet 2022.1.0 install, and Lulesh datasets.
  - Alternatively, the sample Jupyter notebook and the Lulesh datasets are available directly at <https://github.com/llnl/hatchet-tutorial>. This repository is integrated with BinderHub, which will create a local interactive environment for you to run the notebook.
- Following this tutorial, you can substitute your own SPOT/Caliper data files into the example notebook.
- We'll use this material in the hands-on portion of the tutorial.

# Automated Application Performance Analysis: Caliper → SPOT → Hatchet

```
#include <caliper/cali.h>

static inline
void LagrangeElements(Domain&
domain, Index_t numElem)
{
    CALI_CXX_MARK_FUNCTION;
// ...
```

Caliper instrumentation  
in the application



At runtime: Performance  
and Metadata Collection



Web-based Visualization and  
Analysis Tools

Analyze  
caliper datasets  
in Python

SPOT and Hatchet

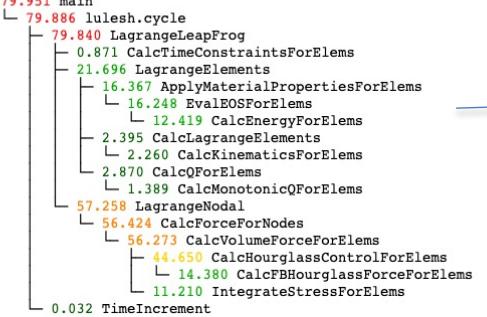
\*Hatchet can analyze other datasets (HPCToolkit, gprof, TAU, Ascent (WIP))

c/o D Boehme

# SPOT Web Interface: Run Table and Jupyter Notebooks

All records selected. Please click on the graph to apply filters.

TABLE    COMPARISON

Figure_of_merit	Problem_size	Threads	Jobsize	launchdate	
1	5491.526855	60	36	1	2019-Jul-16 14:03
	12412.877512	30	4	8	2019-Jul-16 14:04
	5525.153912	60	36	1	2019-Jul-16 14:05
	11339.841229	30			jupyterhub 190716-140428166192 Last Checkpoint: 11/30/2020 (unsaved changes)
	5100.923858	60			
	11197.780437	30			In [4]: # Print the tree representation using the inclusive time metric print(gf.tree(metric_column="time (inc)"))
	5071.148951	60			
	10247.931707	30			
	18468.799766	60			
	18666.324557	60			

Legend (Metric: time (inc))  
71.96 - 79.95  
55.98 - 71.96  
39.99 - 55.98  
24.01 - 39.99  
8.02 - 24.01  
0.03 - 8.02

name User code    Only in left graph    Only in right graph



Buttons bring up Jupyter notebook or specialized analysis views

Jupyter notebook contains Hatchet functions

c/o D Boehme



Lawrence Livermore National Laboratory  
LLNL-PRES-813307





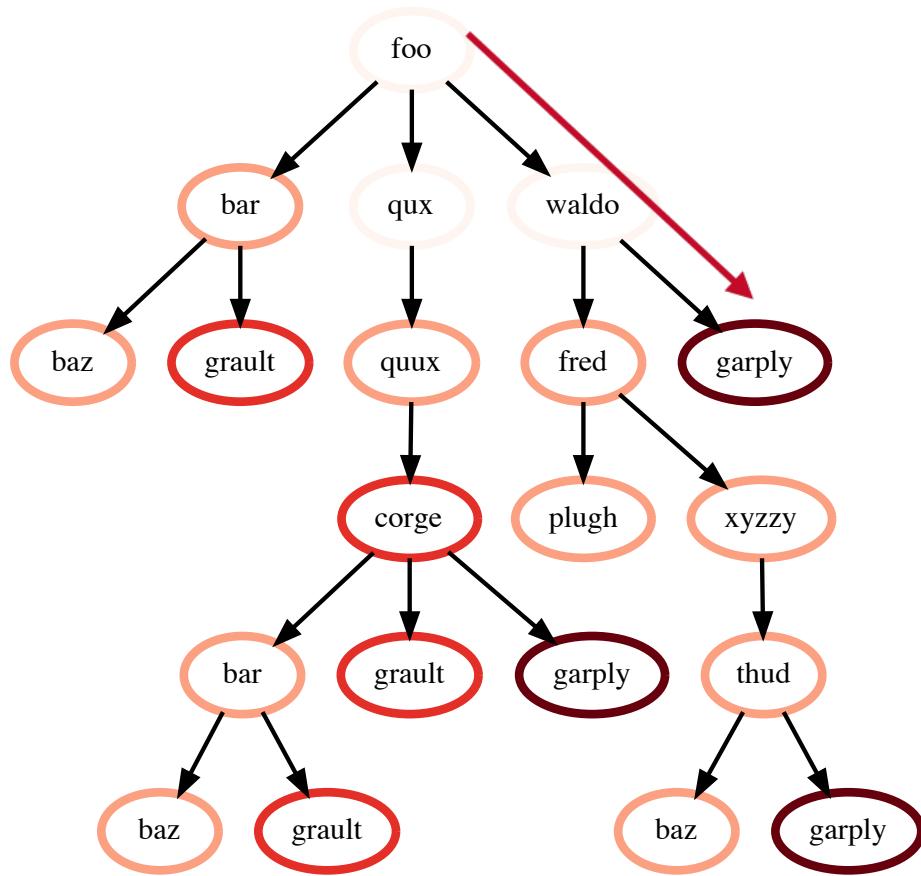
# Hatchet is a performance analysis tool for parallel profiles

- Identify performance bottlenecks to enhance application development
  - Profiling and tracing tools (*e.g.*, Caliper, HPCToolkit, TAU, Score-P, Gprof, Callgrind) provide insights into parts of the code that consume the most time
- Hatchet is an open-source python-based tool for enabling programmatic analysis of structured (or hierarchical) data
- Hatchet can be used to sub-select and focus on a specific region of the data, compare multiple execution profiles, and automate analysis in python scripts



<https://github.com/llnl/hatchet/>

# What do profiling/tracing tools collect?



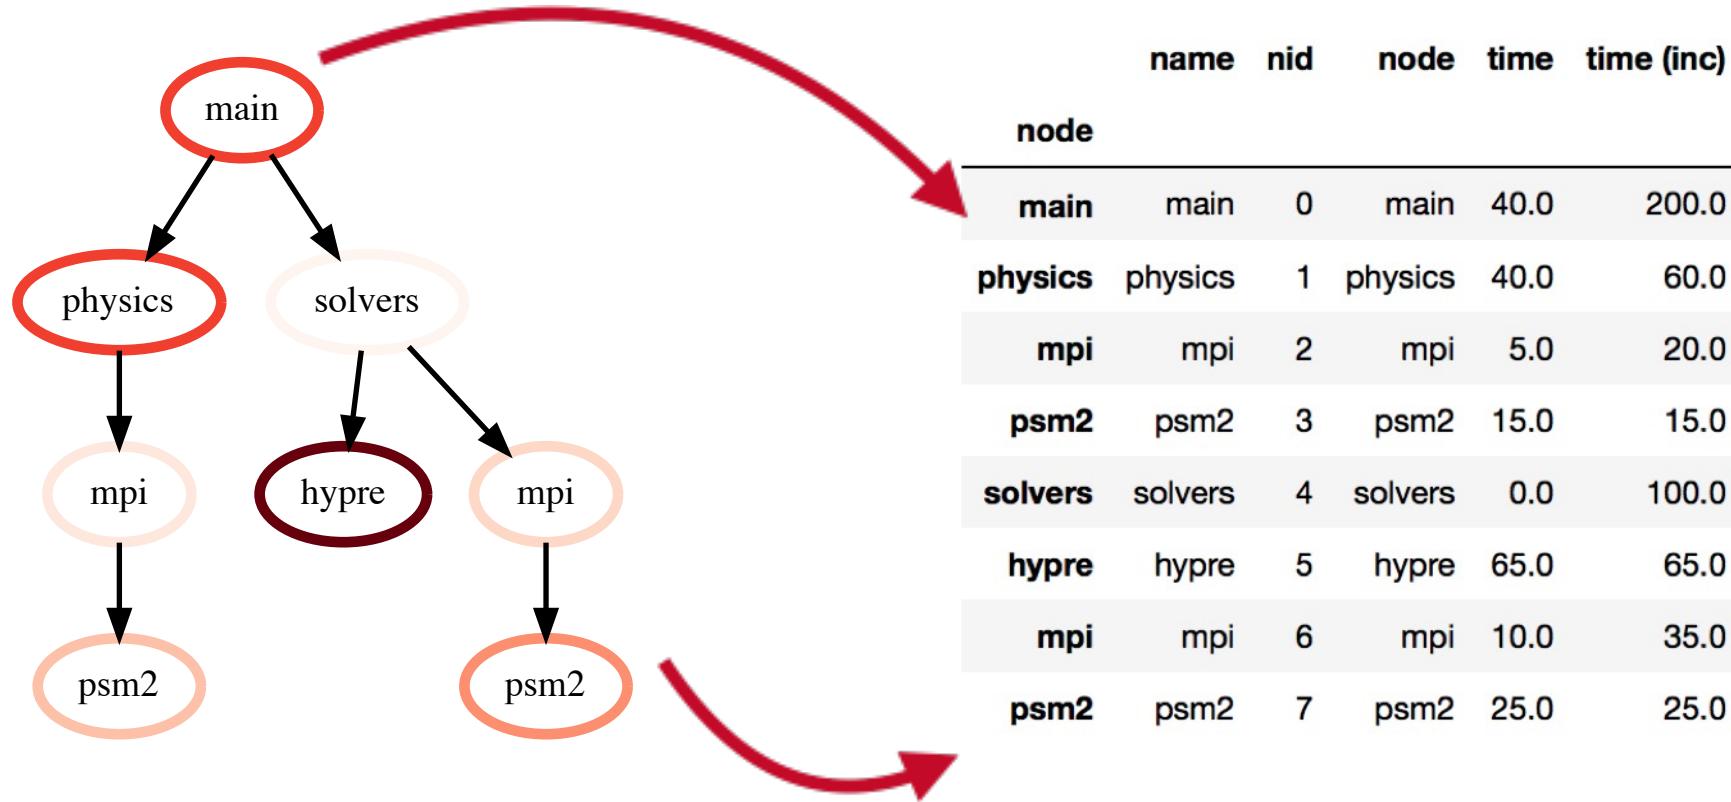
## Calling Context Tree (CCT)

Each node may contain:

- Contextual Info
    - File
    - Line number
    - Function name
    - Callpath
    - Load module
    - Rank ID
    - Thread ID
  - Performance Metrics
    - Time
    - Flops
    - Cache misses

- Hatchet can read profiles from:
  - Caliper
  - HPCToolkit
  - Gprof
  - TAU
  - Ascent (WIP)

# Hatchet's *GraphFrame*: a Graph and a Dataframe



**Graph:** Stores relationships between parents and children

**Pandas Dataframe:** 2D table storing numerical data associated with each node (may be unique per rank, per thread)

# Visualizing Hatchet's GraphFrame components

```
>>> print(gf.tree()) # print graph  
>>> print(gf.dataframe) # print dataframe
```

```
0.000 foo  
|   └ 6.000 bar  
|       |   └ 5.000 baz  
|   └ 0.000 qux  
|       |   └ 5.000 quux  
|           |   └ 10.000 corge  
|           |   └ 15.000 garply  
|               |   └ 1.000 grault  
15.000 waldo  
|   └ 3.000 fred  
|       |   └ 5.000 plugh  
|   └ 15.000 garply
```

Legend (Metric: time)

- 13.50 – 15.00
- 10.50 – 13.50
- 7.50 – 10.50
- 4.50 – 7.50
- 1.50 – 4.50
- 0.00 – 1.50

name User code

◀ Only in left graph

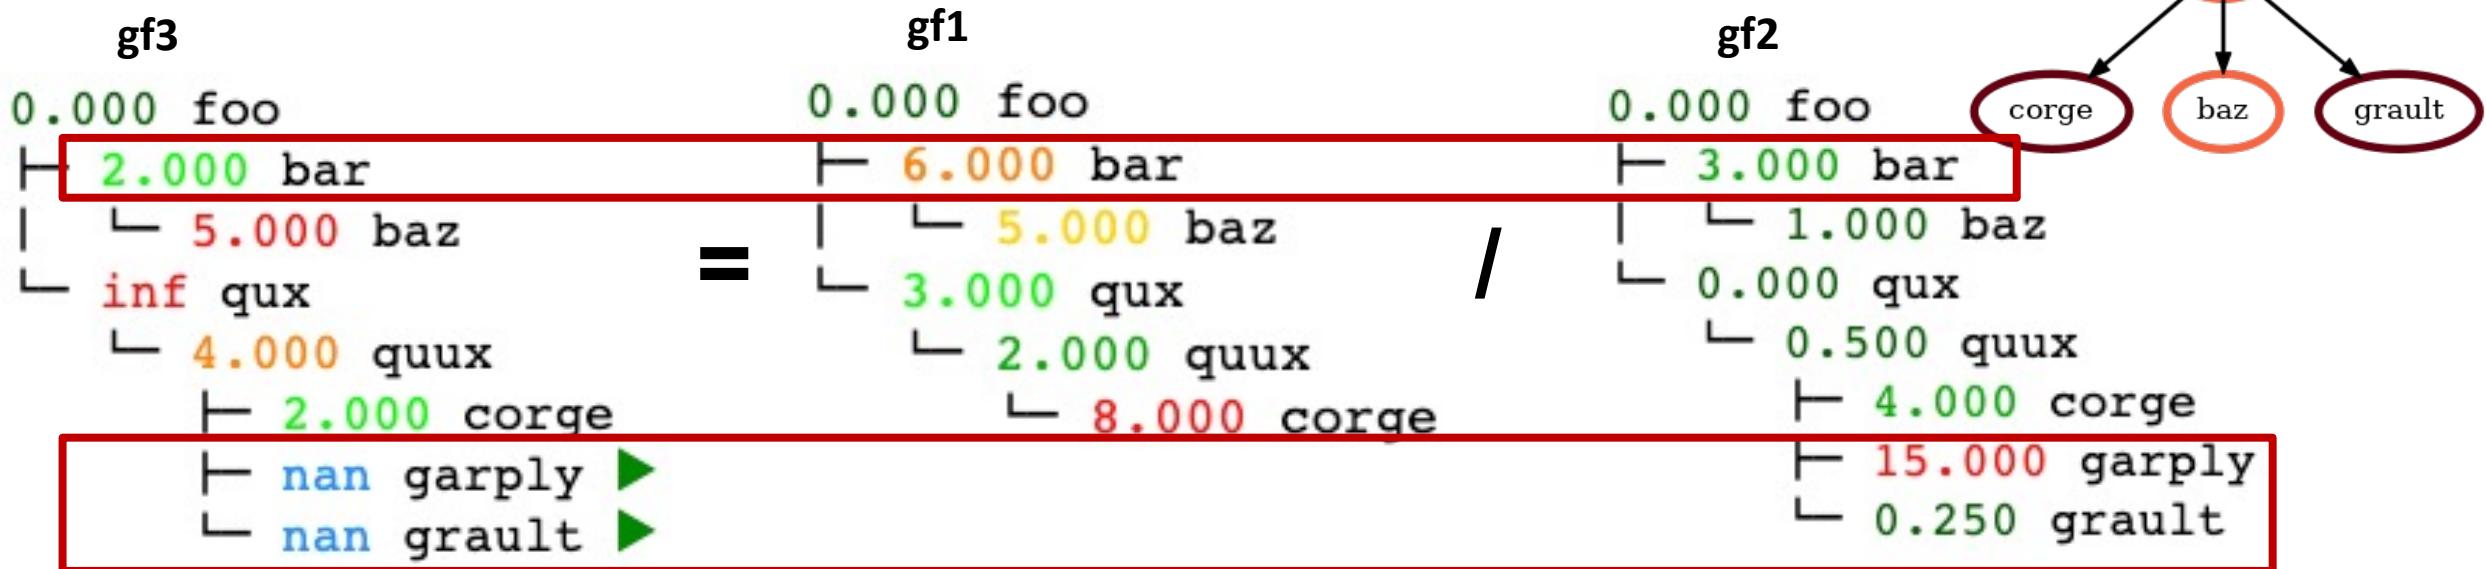
node	name	time	time (inc)
{'name': 'foo'}	foo	0.0	130.0
{'name': 'bar'}	bar	5.0	20.0
{'name': 'baz'}	baz	5.0	5.0
{'name': 'grault'}	grault	10.0	10.0
{'name': 'qux'}	qux	0.0	60.0
{'name': 'quux'}	quux	5.0	60.0
{'name': 'corge'}	corge	10.0	55.0
{'name': 'bar'}	bar	5.0	20.0
{'name': 'baz'}	baz	5.0	5.0
{'name': 'grault'}	grault	10.0	10.0
{'name': 'garply'}	garply	15.0	15.0
{'name': 'grault'}	grault	10.0	10.0

▶ Only in right graph

# Compare GraphFrames using division (or add, subtract, multiply)

```
>>> gf3 = gf1 / gf2 # divide graphframes
```

\*First, unify two trees since  
structure is different



```
>>> gf3 = gf1 + gf2 # add graphframes  
>>> gf3 = gf1 - gf2 # subtract graphframes  
>>> gf3 = gf1 * gf2 # multiply graphframes
```

# Filter the GraphFrame by node metrics in the dataframe

```
>>> filter_func = lambda x: x["time"] > 1 # filter function  
>>> filt_gf = gf.filter(filter_func, squash=True) # apply filter and rewire graph
```

```
0.000 foo  
|   └ 6.000 bar  
|       └ 5.000 baz  
|   └ 0.000 qux  
|       └ 5.000 quux  
|           └ 10.000 corge  
|           └ 15.000 garply  
|               └ 1.000 grault  
└ 15.000 waldo  
    └ 3.000 fred  
        └ 5.000 plugh  
    └ 15.000 garply
```



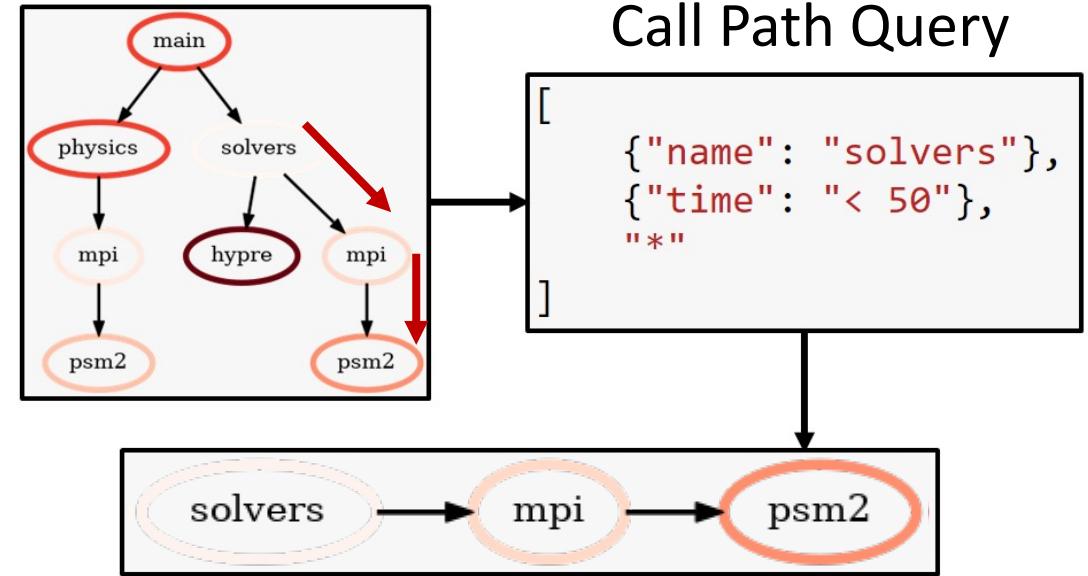
Keep only those nodes with a value greater than 1

```
6.000 bar  
└ 5.000 baz  
5.000 quux  
└ 10.000 corge  
└ 15.000 garply  
15.000 waldo  
└ 3.000 fred  
    └ 5.000 plugh  
    └ 15.000 garply
```

# Filter the GraphFrame using Hatchet's *call path query language*

- Data reduction using *call path* pattern matching

```
# filter using call path query language
query = [
    { "name": "solvers" },
    { "time": "< 50" },
    "*"
]
filt_gf = gf.filter(query, squash=True)
```

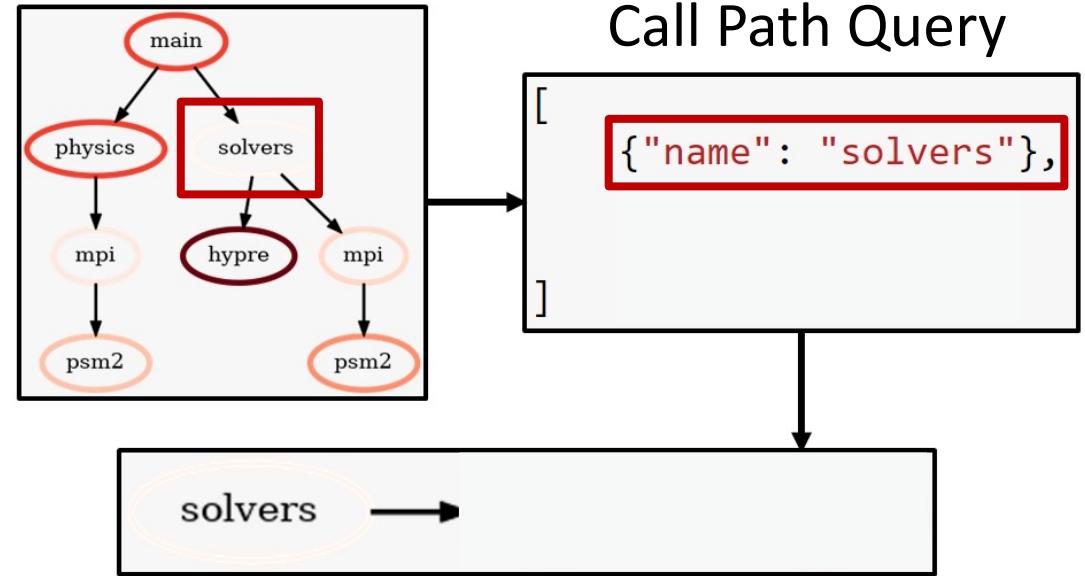


Matches a call path (1) rooted at a node with name “solvers”, (2) followed by a node with a time metric value less than 50, and (3) followed by any number of children nodes.

# Filter the GraphFrame using Hatchet's *call path query language*

- Data reduction using *call path* pattern matching

```
# filter using call path query language
query = [
    { "name": "solvers" },
    { "time": "< 50" },
    "*"
]
filt_gf = gf.filter(query, squash=True)
```

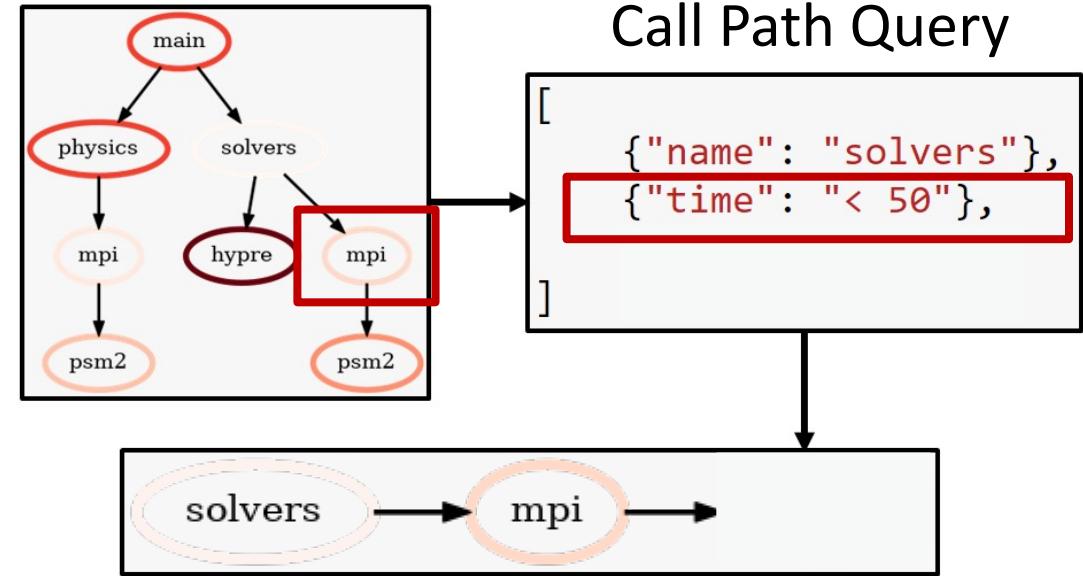


Matches a call path (1) rooted at a node with name “solvers”, (2) followed by a node with a time metric value less than 50, and (3) followed by any number of children nodes.

# Filter the GraphFrame using Hatchet's *call path query language*

- Data reduction using *call path* pattern matching

```
# filter using call path query language
query = [
    { "name": "solvers" },
    { "time": "< 50" },
    "*"
]
filt_gf = gf.filter(query, squash=True)
```

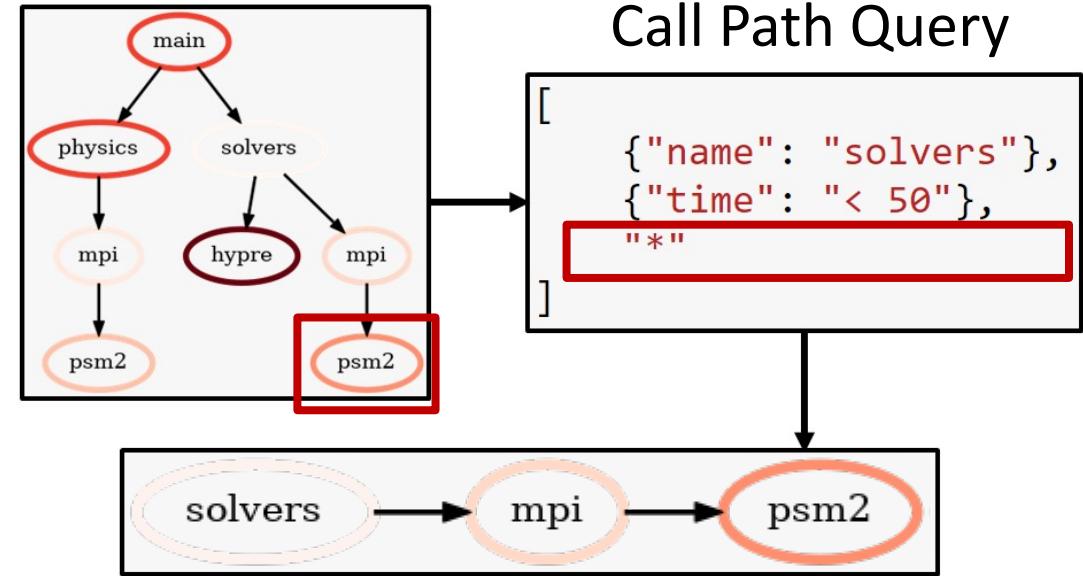


Matches a call path (1) rooted at a node with name “solvers”, (2) followed by a node with a time metric value less than 50, and (3) followed by any number of children nodes.

# Filter the GraphFrame using Hatchet's *call path query language*

- Data reduction using *call path* pattern matching

```
# filter using call path query language
query = [
    { "name": "solvers" },
    { "time": "< 50" },
    "*"
]
filt_gf = gf.filter(query, squash=True)
```



Matches a call path (1) rooted at a node with name “solvers”, (2) followed by a node with a time metric value less than 50, and (3) followed by any number of children nodes.

# How do I load SPOT/Caliper data into Hatchet?

All records selected. Please click on the graph to apply filters.

1. Directory of SPOT/Caliper files

Figure_of_merit	Problem_size	Threads	Jobsize	launchdate
5491.526855	60	36	1	2019-Jul-16 14:03
12412.877512	30	4	8	2019-Jul-16 14:04
5525.153912	60	36	1	2019-Jul-16 14:05
11339.841229	30	4	8	2019-Jul-16 14:06
5100.923858	60	36	1	2019-Jul-16 14:07
11197.780437	30	4	8	2019-Jul-16 14:08
5071.148951	60	36	1	2019-Jul-16 14:09
10247.931707	30	4	8	2019-Jul-16 14:10
18468.799766	60	8	8	2019-Jul-16 14:06
18666.324557	60	8	8	2019-Jul-16 14:13

2. Buttons bring up filled-in Jupyter notebook loading 1 or many SPOT/Caliper files

jupyterhub combo Last Checkpoint: 03/25/2021 (unsaved changes)

In [ ]:

```
import sys
import platform

machine = platform.uname().machine

# Variables replaced by spot.py
input_deploy_dir_str = "/usr/gapps/spot/dev/"
input_db_uri_str = "/usr/gapps/spot/datasets/lulesh_new"
input_run_ids_str = "190716-14510612170.cal"

sys.path.append(input_deploy_dir_str + "/hatchet-venv/" + machine + "/lib/python3.7/site-packages")
sys.path.append(input_deploy_dir_str + "/hatchet/" + machine)
sys.path.append(input_deploy_dir_str + "/spotdb")

import datetime as dt
import hatchet
import spotdb

db = spotdb.connect(input_db_uri_str)
runs = input_run_ids_str.split(',')

gfs = hatchet.GraphFrame.from_spotdb(db, runs)
gf = gfs.pop()

launchdate = dt.datetime.fromtimestamp(int(gf.metadata["launchdate"]))
jobsizes = int(gf.metadata.get("jobsizes", 1))

print("launchdate: {}, jobsizes: {}".format(launchdate, jobsizes))
print(gf.tree())
```

3. Caliper file(s) to explore

4. Connect to SPOT database

5. Hatchet's SPOT database reader loads into Hatchet's GraphFrame object

Documentation: [llnl-hatchet](#)

LLNL-PRES-813307

15

# Hands-On Time!

- The SPOT container includes a sample Jupyter notebook, Hatchet 2022.1.0 install, and Lulesh datasets.
  - Alternatively, the sample Jupyter notebook and the Lulesh datasets are available directly at <https://github.com/llnl/hatchet-tutorial>. This repository is integrated with BinderHub, which will create a local interactive environment for you to run the notebook.
- Following this tutorial, you can substitute your own SPOT/Caliper data files into the example notebook.



- Hop over to Jupyter to run the notebook
- We'll be walking through `hatchet_tutorial_demo.ipynb`

# Review: Topics covered in today's tutorial

- Single graph:
    - Load SPOT/Caliper data file
    - Visualize tree and dataframe
    - Filter and squash tree
  

```
# Read in a SPOT/Caliper file
gf = ht.GraphFrame.from_spotdb(
    <spot-database>,
    <list-of-runs>,
)
```

```
# Print tree visualization
print(gf.tree(metric_column="time (inc)"))
```
- Subtract two trees:
    - Load two SPOT/Caliper data files
    - Compute percent change of two nightly test runs (two different times)
    - Update existing column in dataframe
    - Added new column to dataframe
    - Visualize resulting tree
  

```
# Print dataframe
print(gf.dataframe)
```
- Speedup of two trees:
    - Load two SPOT/Caliper data files
    - Divide two graphs for speedup comparison
    - Visualize resulting tree
    - Generate speedup plot for interesting functions
  

```
# Divide two trees
gf3 = gf2 / gf1
```

```
# Diff two trees
gf3 = (gf2 - gf1) / gf1
```

# Readily available features not covered in today's tutorial

---

- **Add or multiply two graphframes**
- **Insert new column** to dataframe of metrics
  - Scale and offset “time” column by some factor: [https://llnl-hatchet.readthedocs.io/en/latest/advanced\\_examples.html#applying-scalar-operations-to-attributes](https://llnl-hatchet.readthedocs.io/en/latest/advanced_examples.html#applying-scalar-operations-to-attributes)
  - Compute imbalance across MPI ranks within a single application execution: [https://llnl-hatchet.readthedocs.io/en/latest/advanced\\_examples.html#applying-scalar-operations-to-attributes](https://llnl-hatchet.readthedocs.io/en/latest/advanced_examples.html#applying-scalar-operations-to-attributes)
- **Groupby-and-aggregate** nodes by other columns (e.g., function name, file name)
  - `res = gf.groupby_aggregate(["file"], {"time": np.sum})`
- For more details, please visit our User Guide: [https://llnl-hatchet.readthedocs.io/en/latest/user\\_guide.html](https://llnl-hatchet.readthedocs.io/en/latest/user_guide.html)

# Summary

- Hatchet is a performance analysis tool for parallel profiles
- It enables programmatic analysis of hierarchical data from one or multiple execution profiles
- Future Work:
  - Support other profile formats, add a format for outputting GraphFrames to disk
  - Implement a higher-level API for automating performance analysis
- Hatchet      <https://github.com/LLNL/hatchet>  
<https://github.com/LLNL/hatchet-tutorial>
- Caliper      <https://github.com/LLNL/Caliper>
- SPOT      [https://github.com/LLNL/spot2\\_container](https://github.com/LLNL/spot2_container)



Please contact us or submit GitHub issues for Hatchet questions, issues, or feature requests!



**Lawrence Livermore  
National Laboratory**

**Disclaimer**

This document was prepared as an account of work sponsored by an agency of the United States government. Neither the United States government nor Lawrence Livermore National Security, LLC, nor any of their employees makes any warranty, expressed or implied, or assumes any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represents that its use would not infringe privately owned rights. Reference herein to any specific commercial product, process, or service by trade name, trademark, manufacturer, or otherwise does not necessarily constitute or imply its endorsement, recommendation, or favoring by the United States government or Lawrence Livermore National Security, LLC. The views and opinions of authors expressed herein do not necessarily state or reflect those of the United States government or Lawrence Livermore National Security, LLC, and shall not be used for advertising or product endorsement purposes.

# Overview of Hatchet Tutorial Examples

The screenshot shows a Jupyter Notebook interface with the title "jupyter hatchet\_ecp2021\_tutorial\_demo" and a status bar indicating "Last Checkpoint: 22 minutes ago (autosaved)". The notebook has tabs for File, Edit, View, Insert, Cell, Kernel, Help, Trusted, and Python 3. Below the tabs is a toolbar with icons for file operations like Open, Save, Run, and Cell.

**Setup cali-query to extract performance data**

```
In [ ]: import sys  
import os  
import platform  
  
import pandas as pd  
import numpy as np  
from IPython.display import display, HTML  
  
machine = platform.uname().machine  
  
# Add hatchet to PYTHONPATH  
deploy_dir = "/usr/gapps/spot/"  
sys.path.append(deploy_dir + 'hatchet/' + machine)  
import hatchet as ht
```

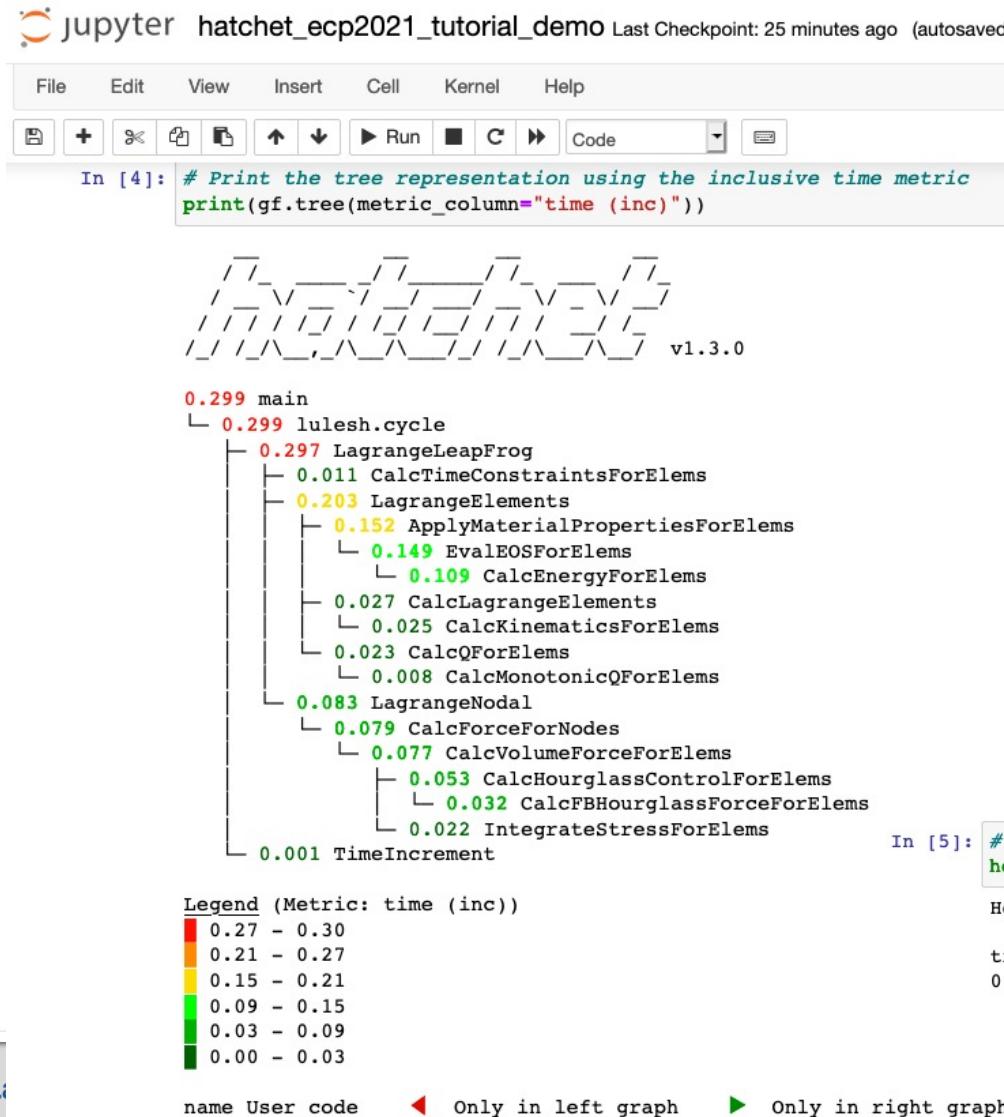
**Load in a single data file and visualize the tree and dataframe**

```
In [ ]: # Path to a Spot/Caliper file  
cali_file = "./hatchet-data/cDPu64825TuLB5ujG_0.cal" # problemsize=10, iter=215, jobsize=1, 2/1/21 8:04  
  
# Read Spot/Caliper file into a Hatchet GraphFrame  
gf = ht.GraphFrame.from_caliper(cali_file, query)  
  
In [ ]: # Print the tree representation using the inclusive time metric  
print(gf.tree(metric_column="time (inc)"))
```

**Caliper file to explore**

**Hatchet's Caliper reader loads into Hatchet's data object called a GraphFrame**

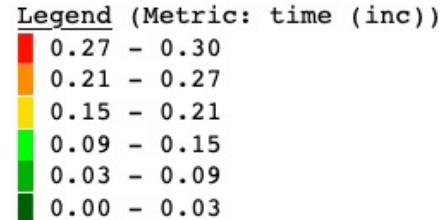
# Visualizing the call graph



- Each node in the tree visualization maps to a function call in the application
- Nodes that are red have high execution time
- Nodes highlighted in grey indicate user functions (code from external libraries are not highlighted)

# Overview of single tree dataset (Lulesh Data)

```
0.299 main
└─ 0.299 lulesh.cycle
   └─ 0.297 LagrangeLeapFrog
      └─ 0.011 CalcTimeConstraintsForElems
         └─ 0.203 LagrangeElements
            └─ 0.152 ApplyMaterialPropertiesForElems
               └─ 0.149 EvalEOSForElems
                  └─ 0.109 CalcEnergyForElems
            └─ 0.027 CalcLagrangeElements
               └─ 0.025 CalcKinematicsForElems
            └─ 0.023 CalcQForElems
               └─ 0.008 CalcMonotonicQForElems
         └─ 0.083 LagrangeNodal
            └─ 0.079 CalcForceForNodes
               └─ 0.077 CalcVolumeForceForElems
                  └─ 0.053 CalcHourglassControlForElems
                     └─ 0.032 CalcFBHourglassForceForElems
                  └─ 0.022 IntegrateStressForElems
   └─ 0.001 TimeIncrement
```



name User code

◀ Only in left graph

▶ Only in right graph

		time (inc)	nid	name
node				
	{'name': 'main', 'type': 'region'}	0.299168	0	main
	{'name': 'lulesh.cycle', 'type': 'region'}	0.298705	1	lulesh.cycle
	{'name': 'LagrangeLeapFrog', 'type': 'region'}	0.297491	3	LagrangeLeapFrog
	{'name': 'CalcTimeConstraintsForElems', 'type': 'region'}	0.010763	18	CalcTimeConstraintsForElems
	{'name': 'LagrangeElements', 'type': 'region'}	0.203010	10	LagrangeElements
	{'name': 'ApplyMaterialPropertiesForElems', 'type': 'region'}	0.151823	15	ApplyMaterialPropertiesForElems
	{'name': 'EvalEOSForElems', 'type': 'region'}	0.149239	16	EvalEOSForElems
	{'name': 'CalcEnergyForElems', 'type': 'region'}	0.108948	17	CalcEnergyForElems
	{'name': 'CalcLagrangeElements', 'type': 'region'}	0.026737	11	CalcLagrangeElements
	{'name': 'CalcKinematicsForElems', 'type': 'region'}	0.025106	12	CalcKinematicsForElems
	{'name': 'CalcQForElems', 'type': 'region'}	0.023439	13	CalcQForElems
	{'name': 'CalcMonotonicQForElems', 'type': 'region'}	0.008134	14	CalcMonotonicQForElems
	{'name': 'LagrangeNodal', 'type': 'region'}	0.083145	4	LagrangeNodal
	{'name': 'CalcForceForNodes', 'type': 'region'}	0.078704	5	CalcForceForNodes
	{'name': 'CalcVolumeForceForElems', 'type': 'region'}	0.076960	6	CalcVolumeForceForElems
	{'name': 'CalcHourglassControlForElems', 'type': 'region'}	0.053464	8	CalcHourglassControlForElems
	{'name': 'CalcFBHourglassForceForElems', 'type': 'region'}	0.032474	9	CalcFBHourglassForceForElems
	{'name': 'IntegrateStressForElems', 'type': 'region'}	0.021547	7	IntegrateStressForElems
	{'name': 'TimeIncrement', 'type': 'region'}	0.000633	2	TimeIncrement

# Filtering a tree

## Original Graph

```
0.299 main
└ 0.299 lulesh.cycle
  └ 0.297 LagrangeLeapFrog
    └ 0.011 CalcTimeConstraintsForElems
      └ 0.203 LagrangeElements
        └ 0.152 ApplyMaterialPropertiesForElems
          └ 0.149 EvalEOSForElems
            └ 0.109 CalcEnergyForElems
          └ 0.027 CalcLagrangeElements
            └ 0.025 CalcKinematicsForElems
          └ 0.023 CalcQForElems
            └ 0.008 CalcMonotonicQForElems
        └ 0.083 LagrangeNodal
          └ 0.079 CalcForceForNodes
            └ 0.077 CalcVolumeForceForElems
              └ 0.053 CalcHourglassControlForElems
                └ 0.032 CalcFBHourglassForceForElems
              └ 0.022 IntegrateStressForElems
    └ 0.001 TimeIncrement
```

Legend (Metric: time (inc))

0.27 - 0.30
0.21 - 0.27
0.15 - 0.21
0.09 - 0.15
0.03 - 0.09
0.00 - 0.03

name User code ◀ Only in left graph ▶ Only in right graph

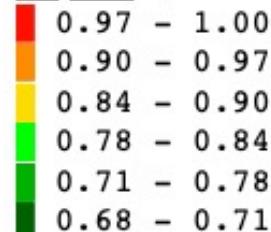
```
# Add new column to the dataframe transforming the inclusive time column to a percentage of the max inclusive time
max_time = gf.dataframe["time (inc)"].max()
gf.dataframe["pct-of-max"] = gf.dataframe["time (inc)"] / max_time
# Filter the tree to contain only nodes consuming at least 60% of max time
filter_func = lambda x: x["pct-of-max"] > 0.6
filtered_squashed_gf = gf.filter(filter_func,
                                  squash=True)
```

Filter graph to keep nodes whose time is greater than 60% of the max time, then rewire graph

## Filtered Graph

```
1.000 main
└ 0.998 lulesh.cycle
  └ 0.994 LagrangeLeapFrog
    └ 0.679 LagrangeElements
```

Legend (Metric: pct-of-max)



name User code ◀ Only in left graph ▶ Only in right graph

# Computing percent change between two trees



# Computing speedup of two trees

0.299 main  
└ 0.299 lulesh.cycle  
  └ 0.297 LagrangeLeapFrog  
    └ 0.011 CalcTimeConstraintsForElems  
      └ 0.203 LagrangeElements  
        └ 0.152 ApplyMaterialPropertiesForElems  
          └ 0.149 EvalEOSForElems  
            └ 0.109 CalcEnergyForElems  
        └ 0.027 CalcLagrangeElements  
        └ 0.025 CalcKinematicsForElems  
      └ 0.023 CalcQForElems  
        └ 0.008 CalcMonotonicQForElems  
    └ 0.083 LagrangeNodal  
      └ 0.079 CalcForceForNodes  
        └ 0.077 CalcVolumeForceForElems  
          └ 0.053 CalcHourglassControlForElems  
            └ 0.032 CalcFBHourglassForceForElems  
        └ 0.022 IntegrateStressForElems  
0.001 TimeIncrement

1 Rank

0.000 MPI\_Allreduce  
0.001 MPI\_Bcast  
0.002 MPI\_Comm\_dup  
0.000 MPI\_Comm\_free  
0.001 MPI\_Comm\_split  
0.000 MPI\_Gather  
0.000 MPI\_Initialized  
3.031 main  
  └ 0.000 MPI\_Barrier  
  └ 0.001 MPI\_Irecv  
  └ 0.000 MPI\_Isend  
  └ 0.000 MPI\_Reduce  
  └ 0.000 MPI\_Wait  
  └ 0.000 MPI\_Waitall  
3.029 lulesh.cycle  
  └ 2.203 LagrangeLeapFrog  
    └ 0.017 CalcTimeConstraintsForElems  
    └ 0.899 LagrangeElements  
      └ 0.420 ApplyMaterialPropertiesForElems  
        └ 0.405 EvalEOSForElems  
          └ 0.265 CalcEnergyForElems  
      └ 0.167 CalcLagrangeElements  
        └ 0.158 CalcKinematicsForElems  
      └ 0.306 CalcQForElems  
        └ 0.045 CalcMonotonicQForElems  
        └ 0.005 MPI\_Irecv  
        └ 0.010 MPI\_Isend  
        └ 0.103 MPI\_Wait  
        └ 0.022 MPI\_Waitall  
1.283 LagrangeNodal  
  └ 1.100 CalcForceForNodes  
    └ 0.756 CalcVolumeForceForElems  
      └ 0.570 CalcHourglassControlForElems  
        └ 0.185 CalcFBHourglassForceForElems  
        └ 0.144 IntegrateStressForElems  
    └ 0.015 MPI\_Irecv  
    └ 0.027 MPI\_Isend  
    └ 0.122 MPI\_Wait  
    └ 0.104 MPI\_Waitall  
    └ 0.008 MPI\_Irecv  
    └ 0.015 MPI\_Isend  
    └ 0.072 MPI\_Wait  
    └ 0.024 MPI\_Waitall  
0.822 TimeIncrement  
  └ 0.820 MPI\_Allreduce

64 Ranks

nan MPI\_Allreduce ►  
nan MPI\_Bcast ►  
nan MPI\_Comm\_dup ►  
nan MPI\_Comm\_free ►  
nan MPI\_Comm\_split ►  
nan MPI\_Gather ►  
nan MPI\_Initialized ►  
0.099 main  
  └ nan MPI\_Barrier ►  
  └ nan MPI\_Irecv ►  
  └ nan MPI\_Isend ►  
  └ nan MPI\_Reduce ►  
  └ nan MPI\_Wait ►  
  └ nan MPI\_Waitall ►  
0.099 lulesh.cycle  
  └ 0.135 LagrangeLeapFrog  
    └ 0.621 CalcTimeConstraintsForElems  
    └ 0.226 LagrangeElements  
      └ 0.361 ApplyMaterialPropertiesForElems  
        └ 0.368 EvalEOSForElems  
          └ 0.411 CalcEnergyForElems  
      └ 0.160 CalcLagrangeElements  
        └ 0.159 CalcKinematicsForElems  
      └ 0.077 CalcQForElems  
        └ 0.181 CalcMonotonicQForElems  
        └ nan MPI\_Irecv ►  
        └ nan MPI\_Isend ►  
        └ nan MPI\_Wait ►  
        └ nan MPI\_Waitall ►  
0.065 LagrangeNodal  
  └ 0.072 CalcForceForNodes  
    └ 0.102 CalcVolumeForceForElems  
      └ 0.094 CalcHourglassControlForElems  
        └ 0.176 CalcFBHourglassForceForElems  
        └ 0.149 IntegrateStressForElems  
      └ nan MPI\_Irecv ►  
      └ nan MPI\_Isend ►  
      └ nan MPI\_Wait ►  
      └ nan MPI\_Waitall ►  
      └ nan MPI\_Irecv ►  
      └ nan MPI\_Isend ►  
      └ nan MPI\_Wait ►  
      └ nan MPI\_Waitall ►  
0.001 TimeIncrement  
  └ nan MPI\_Allreduce ►

Legend (Metric: time (inc))

0.56 - 0.62
0.43 - 0.56
0.31 - 0.43
0.19 - 0.31
0.06 - 0.19
0.00 - 0.06



# Computing speedup of two trees (invert color scheme of result)

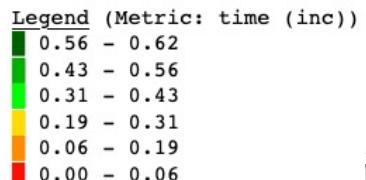
0.299 main  
└ 0.299 lulesh.cycle  
  └ 0.297 LagrangeLeapFrog  
    └ 0.011 CalcTimeConstraintsForElems  
      └ 0.203 LagrangeElements  
        └ 0.152 ApplyMaterialPropertiesForElems  
          └ 0.149 EvalEOSForElems  
            └ 0.109 CalcEnergyForElems  
        └ 0.027 CalcLagrangeElements  
        └ 0.025 CalcKinematicsForElems  
      └ 0.023 CalcQForElems  
        └ 0.008 CalcMonotonicQForElems  
    └ 0.083 LagrangeNodal  
      └ 0.079 CalcForceForNodes  
        └ 0.077 CalcVolumeForceForElems  
          └ 0.053 CalcHourglassControlForElems  
            └ 0.032 CalcFBHourglassForceForElems  
        └ 0.022 IntegrateStressForElems  
0.001 TimeIncrement

1 Rank

0.000 MPI\_Allreduce  
0.001 MPI\_Bcast  
0.002 MPI\_Comm\_dup  
0.000 MPI\_Comm\_free  
0.001 MPI\_Comm\_split  
0.000 MPI\_Gather  
0.000 MPI\_Initialized  
3.031 main  
  └ 0.000 MPI\_Barrier  
  └ 0.001 MPI\_Irecv  
  └ 0.000 MPI\_Isend  
  └ 0.000 MPI\_Reduce  
  └ 0.000 MPI\_Wait  
  └ 0.000 MPI\_Waitall  
3.029 lulesh.cycle  
  └ 2.203 LagrangeLeapFrog  
    └ 0.017 CalcTimeConstraintsForElems  
    └ 0.899 LagrangeElements  
      └ 0.420 ApplyMaterialPropertiesForElems  
        └ 0.405 EvalEOSForElems  
          └ 0.265 CalcEnergyForElems  
        └ 0.167 CalcLagrangeElements  
          └ 0.158 CalcKinematicsForElems  
      └ 0.306 CalcQForElems  
        └ 0.045 CalcMonotonicQForElems  
        └ 0.005 MPI\_Irecv  
        └ 0.010 MPI\_Isend  
        └ 0.103 MPI\_Wait  
        └ 0.022 MPI\_Waitall  
1.283 LagrangeNodal  
  └ 1.100 CalcForceForNodes  
    └ 0.756 CalcVolumeForceForElems  
      └ 0.570 CalcHourglassControlForElems  
        └ 0.185 CalcFBHourglassForceForElems  
        └ 0.144 IntegrateStressForElems  
    └ 0.015 MPI\_Irecv  
    └ 0.027 MPI\_Isend  
    └ 0.122 MPI\_Wait  
    └ 0.104 MPI\_Waitall  
    └ 0.008 MPI\_Irecv  
    └ 0.015 MPI\_Isend  
    └ 0.072 MPI\_Wait  
    └ 0.024 MPI\_Waitall  
0.822 TimeIncrement  
  └ 0.820 MPI\_Allreduce

64 Ranks

nan MPI\_Allreduce ►  
nan MPI\_Bcast ►  
nan MPI\_Comm\_dup ►  
nan MPI\_Comm\_free ►  
nan MPI\_Comm\_split ►  
nan MPI\_Gather ►  
nan MPI\_Initialized ►  
0.099 main  
  └ nan MPI\_Barrier ►  
  └ nan MPI\_Irecv ►  
  └ nan MPI\_Isend ►  
  └ nan MPI\_Reduce ►  
  └ nan MPI\_Wait ►  
  └ nan MPI\_Waitall ►  
0.099 lulesh.cycle  
  └ 0.135 LagrangeLeapFrog  
    └ 0.621 CalcTimeConstraintsForElems  
    └ 0.226 LagrangeElements  
      └ 0.361 ApplyMaterialPropertiesForElems  
        └ 0.368 EvalEOSForElems  
          └ 0.411 CalcEnergyForElems  
        └ 0.160 CalcLagrangeElements  
          └ 0.159 CalcKinematicsForElems  
      └ 0.077 CalcQForElems  
        └ 0.181 CalcMonotonicQForElems  
        └ nan MPI\_Irecv ►  
        └ nan MPI\_Isend ►  
        └ nan MPI\_Wait ►  
        └ nan MPI\_Waitall ►  
0.065 LagrangeNodal  
  └ 0.072 CalcForceForNodes  
    └ 0.102 CalcVolumeForceForElems  
      └ 0.094 CalcHourglassControlForElems  
        └ 0.176 CalcFBHourglassForceForElems  
        └ 0.149 IntegrateStressForElems  
      └ nan MPI\_Irecv ►  
      └ nan MPI\_Isend ►  
      └ nan MPI\_Wait ►  
      └ nan MPI\_Waitall ►  
    └ nan MPI\_Irecv ►  
    └ nan MPI\_Isend ►  
    └ nan MPI\_Wait ►  
    └ nan MPI\_Waitall ►  
0.001 TimeIncrement  
  └ nan MPI\_Allreduce ►



# Generate Lulesh weak scaling plot

